

Introduction

What is R?

R is a free and open source computer program that runs on all major operating systems. R relies primarily on a *command line* interface for data input: instead of interacting with the program by moving your mouse around clicking on different parts of the screen, users enter commands via the keyboard. This will seem to strange to people accustomed to relying on a graphical user interface (GUI) for most of their computing, e.g. via popular programs such as Microsoft Excel or SPSS, yet the approach has a number of benefits, as highlighted by Gary Sherman (2008, p. 283), developer of the popular GIS program QGIS:

With the advent of “modern” GIS software, most people want to point and click their way through life. That’s good, but there is a tremendous amount of flexibility and power waiting for you with the command line. Many times you can do something on the command line in a fraction of the time you can do it with a GUI.

The joy of this, when you get accustomed to it, is that any command is only ever a few keystrokes away, and the order of the commands sent to R can be stored and repeated in scripts, saving even more time in the long-term (more on this in section ...).

Another important attribute of R, related to its command line interface, is that it is a fully fledged *programming language*. Other GIS programs are written in lower level languages such as C++ which are kept at a safe distance from the users by the GUI. In R, by contrast, the user is ‘close to the metal’ in the sense that what he or she inputs is the same as what R sees when it processes the request. This ‘openness’ can seem raw and daunting to beginners, but it is vital to R’s success. Access to R’s source code and openness about how it works has enabled a veritable army of programmers to improve R over time and add an incredible number of extensions to its base capabilities. Consider for a moment that there are now more than 4000 official packages for R, allowing it to tackle almost any computational or numerical problem one could image, and many more that one could not!

Although writing R source code and creating new packages will not appeal to most R users, it inspires confidence to know that there is a strong and highly skilled community of R developers. If there is a useful spatial function that R cannot currently perform, there is a reasonable chance that someone is working on a solution that will become available at a later date. This constant evolution and improvement is a feature of open source software projects not limited to R, but the range and diversity of extensions is certainly one of its strong points. One area where extension of R’s basic capabilities has been particularly successful is the addition of a wide variety of spatial tools.

The rise of R's spatial capabilities

!!! Quick history of R's spatial packages emphasizing current growth and heavy dependence on `sp`.

Mention exciting and recently added packages.

Why R for spatial data visualisation?

Aside from confusion surrounding its one character name - “what kind of a name is R?” [1] and “how can you possibly find resources for R online?” [2] - R may also seem a strange choice for a chapter on *spatial* data visualisation specifically. “I thought R was just for statistics?” and “Why not use a proper GIS package like QGIS?” are valid questions.

The first question arises because R was traditionally conceived - and is still primarily known - as a “statistical programming language” (Bivand and Gebhardt 2000). Although R does have cutting edge statistical capabilities, this definition does not do justice to its power and flexibility. Thus, a more accurate albeit longer definition of R is “an integrated suite of software facilities for data manipulation, calculation and graphical display” (Venables et al. 2013). It is important to consider this wider definition before diving into R: it is a fully fledged programming language meaning that it is highly extensible but also that the same result can often be generated in different ways. This can be confusing.

The second question is based on the premise that all ‘proper’ Geographic Information Systems need to operate in the same way, with primacy allocated to a mapping window and a mouse-driven GUI interface. But when we look back at the history of GIS and its definitions, it becomes clear that R *is* fully fledged GIS, when it is set up correctly. All early GIS programs used a command-line interface; GUIs were only developed later as a way to run commands without needing to remember all the command names (although this is largely overcome by good ‘help’ options and auto-completion). A concise definition of a GIS is “a computerized tool for solving geographic problems” (Longley et al. 2005, p. 16) and R certainly enables this. A more expansive definition of GIS is “a powerful set of tools for collecting, storing, retrieving at will, transforming, and displaying spatial data from the real world for a particular set of purposes” (Burrough and McDonnell, 1998, from Bivand et al. 2013, p. 5); R excels at each of these tasks.

That being said, there are a few major differences between R and conventional GIS programs in terms of spatial data visualisation: R is more suited to creating one-off graphics than exploring spatial data interactively on a map. Conventional GIS packages are better at repeated zooming, panning and spatial sub-setting using custom-drawn polygons than R. Use of the `locator` function allows some interactive selection capabilities in R, but these are limited (Bivand et al. 2013, 3.4). Although interactive maps in R can be created (e.g. using the web interface `shiny`), R should not be seen as a direct replacement of dedicated GIS

programs, especially now that there are myriad free options to try (Sherman 2008). One should use the program which is most appropriate for the task: R can tackle almost any spatial visualisation problem and may be the best option in many cases. In others, however, it may be best used alongside other programs (e.g. Google Earth).

While dedicated GIS programs handle spatial data by default and display the results in a single way, there are various options in R that must be decided by the user. This can be daunting. For example, the user must decide whether to use R's base graphics or a dedicated graphics package such as `ggplot2` for mapping. On the other hand, a major benefit of R is that allows spatial and non-spatial analysis to occur in a *consistent* and *cohesive* framework. Another benefit of R for spatial data visualisation lies in the *reproducibility* of its outputs, a feature that we will be using to great effect in this chapter.

R for Reproducible research

!!! Are all the examples going to be reproducible?

All these components - scripting, stability and the ability to embed 'live' code in documents - make R an excellent tool for transparent research. By using R and carefully documenting what has been done, one ensures that the methods used to reach a certain result can be reproduced by anyone anywhere in the world, provided they have access to the input dataset. The RStudio graphical interface with R encourages good documentation. RStudio enables presentations to be created and professional-quality pdf documents to be produced using the custom file formats `.Rpres` and `.rnw`. In fact, this chapter was written in RMarkdown and converted into a pdf document using only RStudio editor!

R in the wild

Examples of where R has had an important visual impact.

Might be good to mention New York Times etc here as key users of R.

An introductory session

The best way to learn to use a new tool is by using it. The metaphor of craft skills is appropriate here: if you wanted to become skilled at scything, for example, you would not spend your time reading about scythes. The same is true of R: the best way to learn how it works is to 'get your hands dirty' and try it out on your own computer. This introductory session will therefore serve as an introduction to R's unique *syntax*, as well an illustration of how other visualisations presented in this chapter can be reproduced.

R's syntax

Objects

Functions and arguments

Most operations that are performed on objects are done using *functions*. Understanding functions and their various *arguments* is key to manipulating and visualising data in R: the more functions and arguments you know, the more you will be able to do. Functions, in broad terms, are operations that change objects in R from one thing to another. In mathematical language, they *map* sets of numbers onto each other. Arguments are the variables or parameters that are fed into functions to alter their behavior. In terms of R's syntax, arguments are separated by commas within the curved brackets that follow from the function's name. A source of confusion with arguments can be that in some cases they can be inserted directly, whereas in others R needs to be told which argument is being referred to, as illustrated in the code below:

```
seq(from = 0, to = 2, by = 0.5)
```

```
## [1] 0.0 0.5 1.0 1.5 2.0
```

```
seq(0, 2, 0.5)
```

```
## [1] 0.0 0.5 1.0 1.5 2.0
```

```
seq(0, 2, length.out = 6)
```

```
## [1] 0.0 0.4 0.8 1.2 1.6 2.0
```

```
seq(0, 2, 6)
```

```
## [1] 0
```

Before learning about specific functions for spatial analysis and visualisation, it is worth taking some time to think about what a function is and how the arguments passed to it affect how it works. The function `plot` is a good example, because it can take many different input datasets and arguments and produces very different results depending on the arguments it is given. Let's start with a basic example:

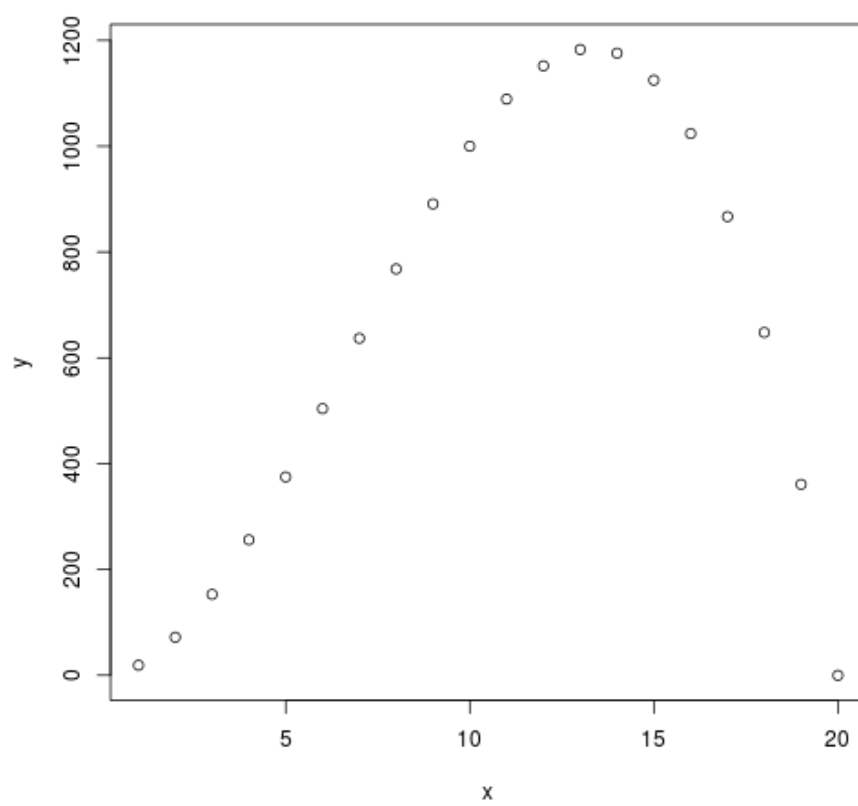


Figure 1: unnamed-chunk-2

```
x <- 1:20
y <- 20 * x^2 - x^3
plot(x, y)
```

In the above code, the function `plot` was given two arguments, `x` and `y` and its default settings are to interpret these as values on a cartesian coordinate system to plot.

Chapter overview

Map Production: Best Practice

Good maps depend on sound analysis and data preparation and can have an enormous impact on the understanding and communication of results. It has never been easier to produce a map. The underlying data required are available in unprecedented volumes and the technological capabilities of transforming them into compelling maps and graphics are increasingly sophisticated and straightforward to use. Data and software, however, only offer the starting points of good spatial data visualisation since they need to be refined and calibrated by the researchers seeking to communicate their findings. In this section we will run through the features of a good map. We will then seek to emulate them with R in Section XX. It is worth noting that not all good maps and graphics contain all the features below – they should simply be seen as suggestions rather than firm principles.

Effective map making is hard process – as Krygier and Wood (XXX) put it “there is a lot to see, think about, and do” (p6). It often comes at the end of a period of intense data analysis and perhaps when the priority is to get a paper finished or results published and can therefore be rushed as a result. The beauty of R (and other scripting languages) is the ability to save code and simply re-run it with different data. Colours, map adornments and other parameters can therefore be quickly applied so it is well worth creating a template script that adheres to best practice.

We have selected `ggplot2` as our package of choice for the bulk of our maps and spatial data visualisations because it has a number of these elements at its core. The “gg” in its slightly odd name stands for “Grammar of Graphics”, which is a set of rules developed by Leland Wilkinson (2005) in a book of the same name. Grammar in the context of graphics works in much the same way as it does in language- it provides a structure. The structure is informed by both human perception and also mathematics to ensure that the resulting visualisations are both technically sound and comprehensible. Through creating `ggplot2`, Hadley Wickham, implemented these rules as well as developing ways in which plots can be built up in layers (see Wickham, 2010). This layering component is especially

useful in the context of spatial data since it is conceptually the same as map layers in Geographical Information Systems (GIS).

First load the libraries required for this section:

```
library(rgdal)

## Loading required package: sp
## rgdal: version: 0.8-14, (SVN revision 496)
## Geospatial Data Abstraction Library extensions to R successfully loaded
## Loaded GDAL runtime: GDAL 1.9.0, released 2011/12/29
## Path to GDAL shared files: /usr/share/gdal/1.9
## Loaded PROJ.4 runtime: Rel. 4.7.1, 23 September 2009, [PJ_VERSION: 470]
## Path to PROJ.4 shared files: (autodetected)
```

```
library(ggplot2)
library(gridExtra)
```

```
## Loading required package: grid
```

You will also need create a folder and then set it as your working directory. Below we assume the name is `Uname`, and the folder is saved as `sdvwR` in the Desktop in Windows.

```
setwd("c:/Users/Uname/Desktop/sdvwR")
```

For this section we are going to use a map of the world to demonstrate some of the cartographic principles discussed. A world map is available from the Natural Earth website. The code below will download this and save it to your working directory.

```
download.file(url = "http://www.naturalearthdata.com/http://www.naturalearthdata.com/download/ne_110m_admin_0_countries.zip", "auto")
unzip("ne_110m_admin_0_countries.zip", exdir = "data/") # unzip to data folder
file.remove("ne_110m_admin_0_countries.zip") # remove zip file

## [1] TRUE
```

Once downloaded we can then load the data into the R console. We have just downloaded a shapefile, which as Section XX explains, is not handled as a “standard” data object in R.

```
wrld <- readOGR("data/", "ne_110m_admin_0_countries")
```

```
## OGR data source with driver: ESRI Shapefile
## Source: "data/", layer: "ne_110m_admin_0_countries"
## with 177 features and 63 fields
## Feature type: wkbPolygon with 2 dimensions

plot(wrld)
```



Figure 2: Initial plot of world boundaries

To see the first ten rows of attribute information associated with each of the country boundaries type the following

```
head(wrld@data)
```

You can see there are a lot of columns associated with this file. Although we will keep all of the them, we are only really interested in the population estimate

("pop_est") field. Before progressing it is worth reprojecting the data in order that the population data can be seen better. The coordinate reference system of the wrld shapefile is currently WGS84. This is the common latitude and longitude format that all spatial software packages understand. From a cartographic perspective the standard plots of this projection, of the kind produced above, are not suitable since they distort the shapes of those countries further from the equator. Instead the Robinson projection provides a good compromise between areal distortion and shape preservation. We therefore project it as follows.

```
library(geosphere)
wrld.rob <- spTransform(wrld, CRS("+proj=robin"))
plot(wrld.rob)
```

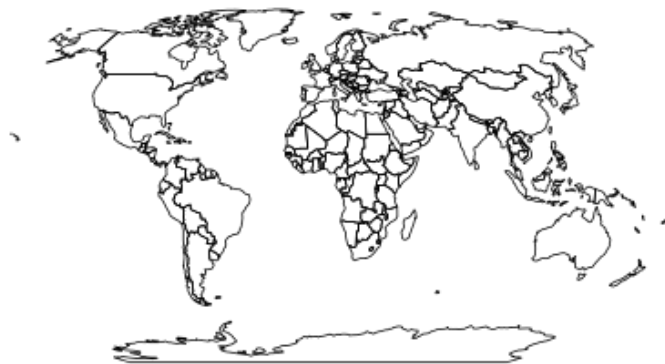


Figure 3: unnamed-chunk-5

“ESRI: 54030” is the reference code of the Robinson projection in the database of projections that R downloads with the `rgdal` package. You will have spotted from the plot that the countries in the world map are much better proportioned.

We now need to “fortify” this spatial data to convert it into a format that `ggplot2` understands, we also use “merge” to re-attach the attribute data that is lost in the fortify operation.

```
# fortify requires rgeos or maptools packages - have we already loaded it?
```

```
# !!!
```

```
wrld.rob.f <- fortify(wrld.rob, region = "sov_a3")
```

```
## Loading required package: rgeos
```

```
## rgeos version: 0.3-2, (SVN revision 413M)
```

```
## GEOS runtime version: 3.3.3-CAPI-1.7.4
```

```
## Polygon checking: TRUE
```

```
wrld.pop.f <- merge(wrld.rob.f, wrld.rob@data, by.x = "id", by.y = "sov_a3")
```

```
# continuous colour ramp
```

```
map <- ggplot(wrld.pop.f, aes(long, lat, group = group, fill = pop_est)) + geom_polygon() +  
  coord_equal() + labs(x = "Longitude", y = "Latitude", fill = "World Population") +  
  ggtitle("World Population")
```

```
# better colours with more breaks- to finish
```

```
map + scale_fill_continuous(breaks = c(10~c(8, 9)))
```

```
# categorical variables
```

Conforming to colour conventions

Colour has an enormous impact on how people will perceive your graphic. “Readers” of a map come to it with a range of pre-conceptions about how the world looks. If the map’s purpose is to clearly communicate data then it is often advisable to conform to conventions so as not to disorientate readers to ensure they can focus on the key messages contained in the data. A good example of this is the use of blue for bodies of water and green for landmass. The code example below generates two plots with our `wrld.pop.f` object. The first colours the land blue and the sea (in this case the background to the map) green and the second is more conventional. We use the “grid.arrange” function from the “gridExtra” package to display the maps side by side.

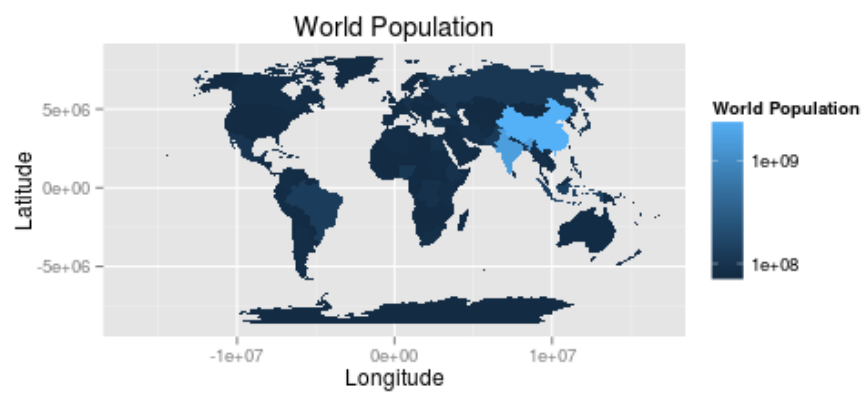


Figure 4: unnamed-chunk-7

```
map2 <- ggplot(wrld.pop.f, aes(long, lat, group = group)) + coord_equal()

blue <- map2 + geom_polygon(fill = "light blue") + theme(panel.background = element_rect(fill = "white", stroke = "black", strokeWidth = 1))

green <- map2 + geom_polygon(fill = "dark green") + theme(panel.background = element_rect(fill = "white", stroke = "black", strokeWidth = 1))

grid.arrange(blue, green, ncol = 2)
```

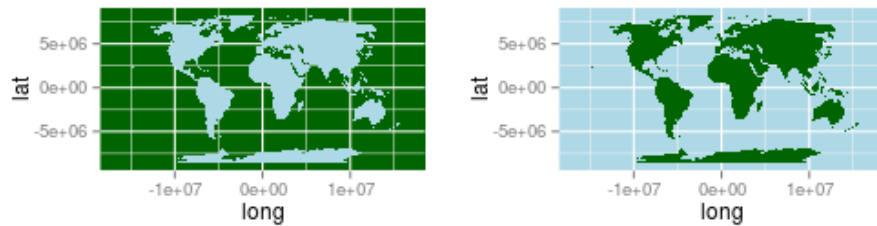


Figure 5: unnamed-chunk-8

Experimenting with line colour and line widths

In addition to conforming to colour conventions, line colour and width offer important parameters, which are often overlooked tools for increasing the leg-

ibility of a graphic. As the code below demonstrates, it is possible to adjust line colour through using the “colour” parameter and the line width using the “lwd” parameter. The impact of different line widths will vary depending on your screen size and resolution. If you save the plot to pdf (or an image) then the size at which you do this will also affect the line widths.

```
map3 <- map2 + theme(panel.background = element_rect(fill = "light blue"))

yellow <- map3 + geom_polygon(fill = "dark green", colour = "yellow")

black <- map3 + geom_polygon(fill = "dark green", colour = "black")

thin <- map3 + geom_polygon(fill = "dark green", colour = "black", lwd = 0.1)

thick <- map3 + geom_polygon(fill = "dark green", colour = "black", lwd = 1.5)

grid.arrange(yellow, black, thick, thin, ncol = 2)
```

There are other parameters such as layer transparency that can be applied to all aspects of the plot - both points, lines and polygons - that we will reference in later examples in this chapter.

Map Adornments and Annotations

Map adornments and annotations are essential to orientate the viewer and provide context; they include graticules, north arrows, scale bars and data attribution. Not all are required on a single map, indeed it is often best that they are used sparingly to avoid unnecessary clutter (Monkhouse and Wilkinson, 1971). Unfortunately it is not always as straightforward to add these in R, and perhaps less so using the ggplot2 paradigm, when compared to a conventional GIS. Here we will outline the ways in which annotations can be added.

!!!! In the maps created so far, we have defined the *aesthetics* of the map in the foundation function `ggplot`. The result of this is that all subsequent layers are expected to have the same variables and essentially contain data with the same dimensions as original dataset. But what if we want to add a new layer from a completely different dataset. To do this, we must not add any arguments to the `ggplot` function, only adding data sources one layer at a time:

North arrow

```
ggplot() + geom_polygon(data = wrld.pop.f, aes(long, lat, group = group, fill = pop_est)) +
  geom_line(aes(x = c(-160, -160), y = c(0, 25)), arrow = arrow())
```

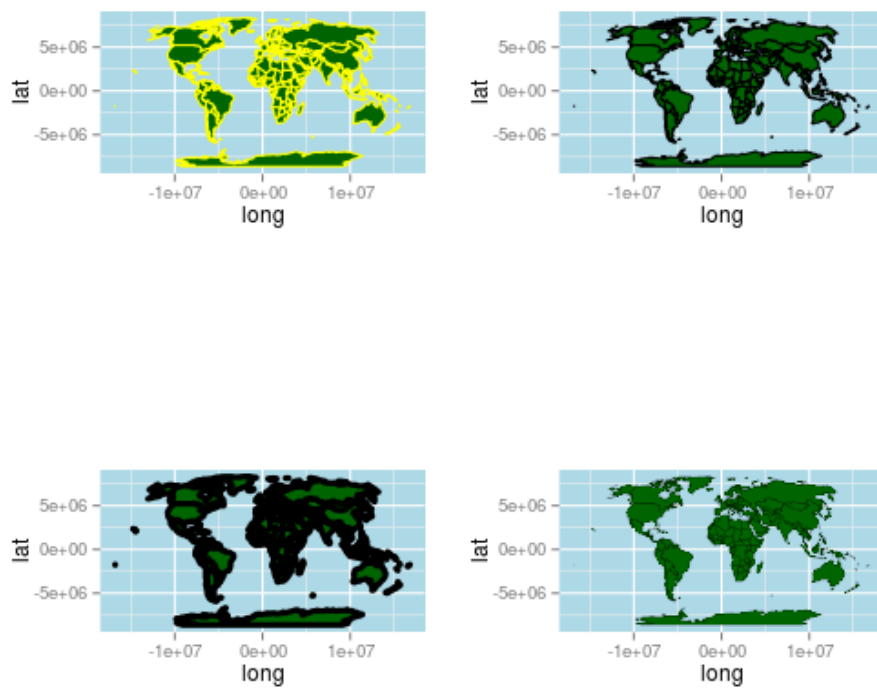


Figure 6: unnamed-chunk-9

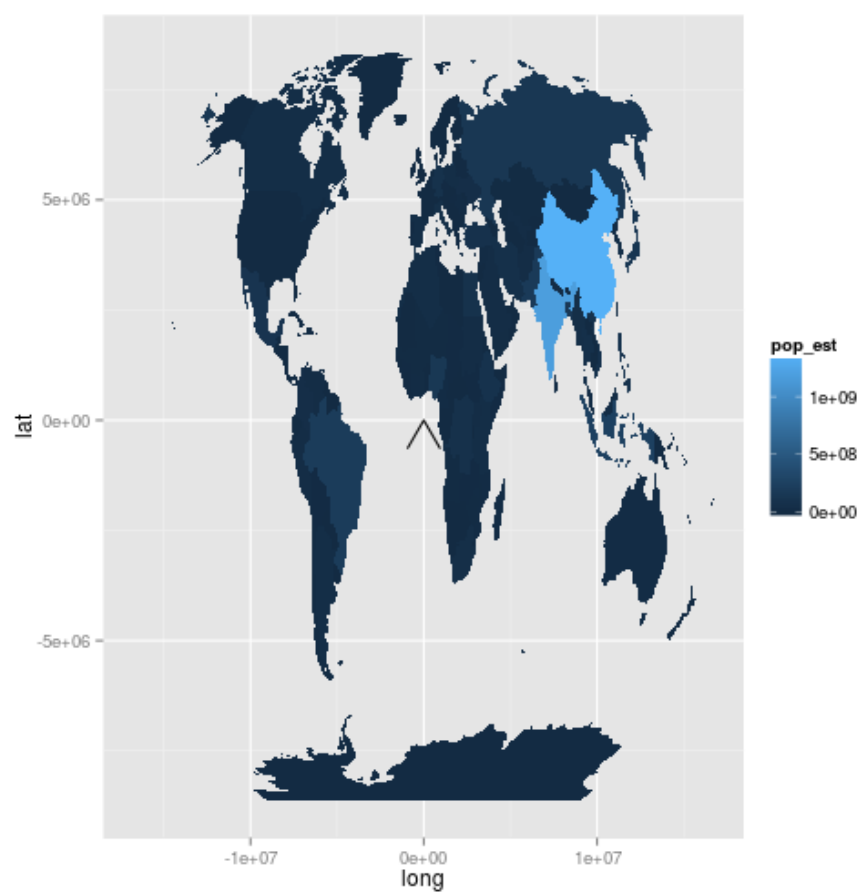


Figure 7: unnamed-chunk-10

```

# scale bar- found this function

hscale_segment = function(breaks, ...) {
  y = unique(breaks$y)
  stopifnot(length(y) == 1)
  dx = max(breaks$x) - min(breaks$x)
  dy = 1/30 * dx
  hscale = data.frame(ix = min(breaks$x), iy = y, jx = max(breaks$x), jy = y)
  vticks = data.frame(ix = breaks$x, iy = (y - dy), jx = breaks$x, jy = (y +
    dy))
  df = rbind(hscale, vticks)
  return(geom_segment(data = df, aes(x = ix, xend = jx, y = iy, yend = jy),
    ...))
}

hscale_text = function(breaks, ...) {
  dx = max(breaks$x) - min(breaks$x)
  dy = 2/30 * dx
  breaks$y = breaks$y + dy
  return(geom_text(data = breaks, aes(x = x, y = y, label = label), hjust = 0.5,
    vjust = 0, ...))
}

```

There is an almost infinite number of different combinations of the above parameters so take inspiration from maps and graphics you have seen and liked. The process is an iterative one, it will take multiple attempts to get right. Show your map to friends and colleagues- all will have an opinion but don't be afraid to stand by the decisions you have taken.

Consistency- across papers.

R and Spatial Data

Spatial Data in R

In any data analysis project, spatial or otherwise, it is important to have a strong understanding of the dataset before progressing. This section will therefore begin with a description of the input data used in this section. We will see how data can be loaded into R and exported to other formats, before going into more detail about the underlying structure of spatial data in R: how it 'sees' spatial data is quite unique.

Loading spatial data in R

In most situations, the starting point of spatial analysis tasks is loading in pre-existing datasets. These may originate from government agencies, remote sensing devices or ‘volunteered geographical information’ from GPS devices, online databases such as Open Street Map or geo-tagged social media (Goodchild 2007). The diversity of geographical data formats is large.

R is able to import a very wide range of spatial data formats thanks to its interface with the Geospatial Data Abstraction Library (GDAL), which is enabled by loading the package `rgdal` into R. Below we will load data from two spatial data formats: GPS eXchange (`.gpx`) and an ESRI Shapefile (consisting of at least files with `.shp`, `.shx` and `.dbf` extensions).

`readOGR` is in fact capable of loading dozens more file formats, so the focus is on the *method* rather than the specific formats. The ‘take home message’ is that the `readOGR` function is capable of loading most common spatial file formats, but behaves differently depending on file type. Let’s start with a `.gpx` file, a tracklog recording a bicycle ride from Sheffield to Wakefield which was uploaded Open Street Map. [!!! more detail?]

```
# download.file('http://www.openstreetmap.org/trace/1619756/data', destfile
# = 'data/gps-trace.gpx')
library(rgdal) # load the gdal package

## Loading required package: sp
## rgdal: version: 0.8-11, (SVN revision 479M)
## Geospatial Data Abstraction Library extensions to R successfully loaded
## Loaded GDAL runtime: GDAL 1.9.2, released 2012/10/08
## Path to GDAL shared files: /usr/share/gdal
## Loaded PROJ.4 runtime: Rel. 4.8.0, 6 March 2012, [PJ_VERSION: 480]
## Path to PROJ.4 shared files: (autodetected)

ogrListLayers(dsn = "data/gps-trace.gpx") # which layers are available?

## [1] "waypoints"      "routes"         "tracks"         "route_points"
## [5] "track_points"

shf2lds <- readOGR(dsn = "data/gps-trace.gpx", layer = "tracks") # load track

## OGR data source with driver: GPX
## Source: "data/gps-trace.gpx", layer: "tracks"
## with 1 features and 12 fields
## Feature type: wkbMultiLineString with 2 dimensions
```

```

plot(shf2lds)
shf2lds.p <- readOGR(dsn = "data/gps-trace.gpx", layer = "track_points") # load points

## OGR data source with driver: GPX
## Source: "data/gps-trace.gpx", layer: "track_points"
## with 6085 features and 26 fields
## Feature type: wkbPoint with 2 dimensions

points(shf2lds.p[seq(1, 3000, 100), ])

```



Figure 8: Leeds to Sheffield GPS data

There is a lot going on in the preceding 7 lines of code, including functions that you are unlikely to have encountered before. Let us think about what has happened, line-by-line.

First, we used R to *download* a file from the internet, using the function `download.file`. The two essential arguments of this function are `url` (we could have typed `url =` before the link) and `destfile` (which means destination file). As with any function, more optional arguments can be viewed by typing `?download.file`.

When `rgdal` has successfully loaded, the next task is not to import the file directly, but to find out which *layers* are available to import, with the function `ogrListLayers`. The output from this command tells us that various layers are available, including `tracks` and `track_points`, which we subsequently load using `readOGR`. The basic `plot` function is used to plot the newly imported objects, ensuring they make sense. In the second `plot` function, we take a subset of the object (see section ... for more on this).

As stated in the help documentation (accessed by entering `?readOGR`), the `dsn` = argument is interpreted differently depending on the type of file used. In the above example, the filename was the data source name. To load Shapefiles, by contrast, the *folder* containing the data is used:

```
lnd <- readOGR(dsn = "data/", "london_sport")
```

Here, the data is assumed to reside in a folder entitled `data` which in R's current working directory (remember to check this using `getwd()`). If the files were stored in the working directory, one would use `dsn = "."` instead. Again, it may be wise to plot the data that results, to ensure that it has worked correctly. Now that the data has been loaded into R's own `sp` format, try interrogating and plotting it, using functions such as `summary` and `plot`.

The size of spatial datasets in R

Any data that has been read into R's *workspace*, which constitutes all objects that can be accessed by name and can be listed using the `ls()` function, can be saved in R's own data storage file type, `.RData`. Spatial datasets can get quite large and this can cause problems on computers by consuming all available random access memory (RAM) or hard disk space available to the computer. It is therefore wise to understand roughly how large spatial objects are; this will also provide insight into how long certain functions will take to run.

In the absence of prior knowledge, which of the two objects loaded in the previous section would be expected to take up more memory. One could hypothesise that the London boroughs represented by the object `lnd` would be larger, but how much larger? We could simply look at the size of the associated files, but R also provides a function (`object.size`) for discovering how large objects loaded into its workspace are:

```
object.size(shf2lds)
```

```
object.size(lnd)
```

Surprisingly, the GPS data is larger. To see why, we can find out how many *vertices* (points connected by lines) are contained in each dataset:

```
##      [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

```
## [1] 1102
```

```
## [1] 1
```

```
## [1] 6085
```

Without worrying, for now, about how these vertex counts were performed, it is clear that the GPS data has almost 6 times the number of vertices as does the London data, explaining its larger size. Yet when plotted, the GPS data does not seem more detailed, implying that some of the vertices in the object are not needed for visualisation at the scale of the objects *bounding box*.

Simplifying geometries

The wastefulness of the GPS data for visualisation (the full dataset may be useful for other types of analysis) raises the question following question: can the object be simplified such that its key features remain while substantially reducing its size? The answer is yes. In the code below, we harness the power of the `rgeos` package and its `gSimplify` function to simplify spatial R objects (the code can also be used to simplify polygon geometries):

```
library(rgeos)

## rgeos version: 0.3-2, (SVN revision 413M)
## GEOS runtime version: 3.3.9-CAPI-1.7.9
## Polygon checking: TRUE

shf2lds.simple <- gSimplify(shf2lds, tol = 0.001)
(object.size(shf2lds.simple)/object.size(shf2lds))[1]

## [1] 0.04608

plot(shf2lds.simple)
plot(shf2lds, col = "red", add = T)
```

In the above block of code, `gSimplify` is given the object `shf2lds` and the `tol` argument, short for “tolerance”, is set at 0.001 (much larger values may be needed, for data that use is *projected* - does not use latitude and longitude). Comparison between the sizes of the simplified object and the original shows that the new object is less than 3% of its original size. Try plotting the original and simplified tracks on your computer: when visualised using the `plot` function, it becomes clear that the object `shf2lds.simple` retains the overall shape of the line and is virtually indistinguishable from the original object.

This example is rather contrived because even the larger object `shf2lds` is only 0.107 Mb, negligible compared with the gigabytes of RAM available to modern computers. However, it underlines a wider point: for *visualisation* purposes at small spatial scales (i.e. covering a large area of the Earth on a small map), the *geometries* associated with spatial data can often be simplified to reduce processing time and usage of RAM. The other advantage of simplification is that it reduces the size occupied by spatial datasets when they are saved.

Saving and exporting spatial objects

The structure of spatial data in R

Spatial* data

Points

Lines

Polygons

Grids and raster data

‘Flattening’ data with `fortify`

The main spatial packages

`sp`

`rgdal`

`rgeos`

Maps with `ggplot2`

Adding base maps with `ggmap`

Manipulating spatial data

Coordinate reference systems and transformations

Attribute joins

Spatial joins

A spatial join, like attribute joins, is used to transfer information from one dataset to another. There is a clearly defined direction to spatial joins, with the *target layer* receiving information from another spatial layer based on the proximity of elements from both layers to each other. There are three broad types of spatial join: one-to-one, many-to-one and one-to-many. We will focus only the former two as the third type is rarely used.

One-to-one spatial joins are by far the easiest to understand and compute because they simply involve the transfer of attributes in one layer to another, based on location. A one-to-one join is depicted in figure x below.

Many-to-one spatial joins involve taking a spatial layer with many elements and allocating the attributes associated with these elements to relatively few elements in the target spatial layer. A common type of many-to-one spatial join is the allocation of data collected at many point sources unevenly scattered over space to polygons representing administrative boundaries, as represented in Fig. x.



Figure 9: Illustration of a one-to-one spatial join

```

lnd.stations <- readOGR("data/", "lnd-stns", p4s = "+init=epsg:27700")

## OGR data source with driver: ESRI Shapefile
## Source: "data/", layer: "lnd-stns"
## with 2532 features and 6 fields
## Feature type: wkbPoint with 2 dimensions

plot(lnd)
plot(lnd.stations[round(runif(n = 500, min = 1, max = nrow(lnd.stations))),
], add = T)

```

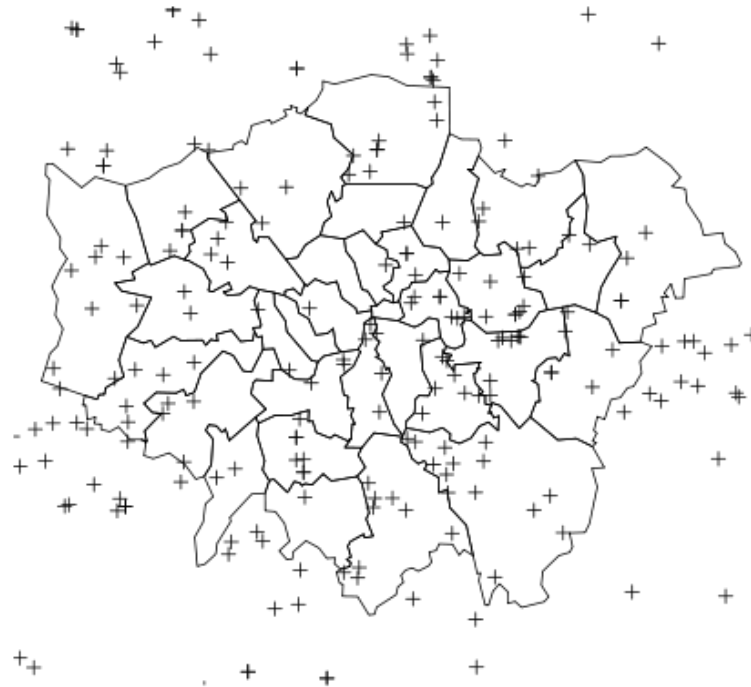


Figure 10: Input data for a spatial join

The above code reads in a `SpatialPointsDataFrame` consisting of 2532 transport nodes in and surrounding London and then plots a random sample of 500 of

these over the previously loaded borough level administrative boundaries. The reason for plotting a sample of the points rather than all of them is that the boundary data becomes difficult to see if all of the points are plotted. It is also useful to see and practice sampling techniques in practice; try to plot only the first 500 points, rather than a random selection, and describe the difference.

The most obvious issue with the point data from the perspective of a spatial join with the borough data is that many of the points in the dataset are in fact located outside the region of interest. Thus, the first stage in the analysis is to filter the point data such that only those that lie within London's administrative zones are selected. This in itself is a kind of spatial join, and can be accomplished with the following code.

```
proj4string(lnd) <- proj4string(lnd.stations)

## Warning: A new CRS was assigned to an object with an existing CRS:
## +proj=tmerc +lat_0=49 +lon_0=-2 +k=0.9996012717 +x_0=400000 +y_0=-100000 +ellps=airy +units=m +no_defs
## without reprojecting.
## For reprojection, use function spTransform in package rgdal

lnd.stations <- lnd.stations[lnd, ] # select only points within lnd
plot(lnd.stations) # check the result
```

The station points now clearly follow the form of the `lnd` shape, indicating that the procedure worked. Let's review the code that allowed this to happen: the first line ensured that the CRS associated with each layer is *exactly* the same: this step should not be required in most cases, but it is worth knowing about. Of course, if the coordinate systems are *actually* different in each layer, the function `spTransform` will be needed to make them compatible. This procedure is discussed in section !!! In this case, only the name was slightly different hence direct alteration of the CRS name via the function `proj4string`.

The second line of code is where the magic happens and the brilliance of R's `sp` package becomes clear: all that was needed was to place another spatial object in the row index of the points (`[lnd,]`) and R automatically understood that a subset based on location should be produced. This line of code is an example of R's 'terseness' - only a single line of code is needed to perform what is in fact quite a complex operation.

Spatial aggregation

Now that only stations which *intersect* with the `lnd` polygon have been selected, the next stage is to extract information about the points within each zone. This many-to-one spatial join is also known as *spatial aggregation*. To do this there

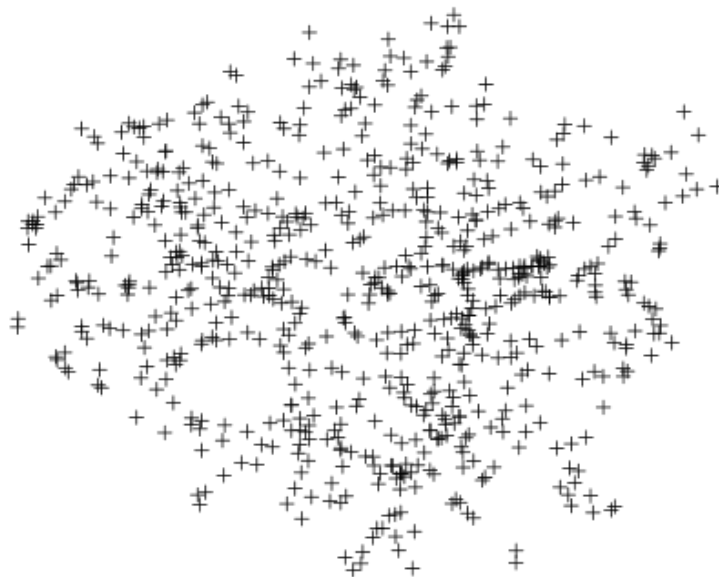


Figure 11: A spatial subset of the points

are a couple of approaches: one using the `sp` package and the other using `rgeos` (see Bivand et al. 2013, 5.3).

As with the *spatial subset* method described above, the developers of R have been very clever in their implementation of spatial aggregations methods. To minimise typing and ensure consistency with R's base functions, `sp` extends the capabilities of the `aggregate` function to automatically detect whether the user is asking for a spatial or a non-spatial aggregation (they are, in essence, the same thing - we recommend learning about the non-spatial use of `aggregate` in R for comparison).

Continuing with the example of station points in London polygons, let us use the spatial extension of `aggregate` to count how many points are in each borough:

```
lndStC <- aggregate(lnd.stations, by = lnd, FUN = length)
summary(lndStC)
plot(lndStC)
```

As with the spatial subset function, the above code is extremely terse. The `aggregate` function here does three things: 1) identifies which stations are in which London borough; 2) uses this information to perform a function on the output, in this case `length`, which simply means “count” in this context; and 3) creates a new spatial object equivalent to `lnd` but with updated attribute data to reflect the results of the spatial aggregation. The results, with a legend and colours added, are presented in Fig !!! below.

As with any spatial attribute data stored as an `sp` object, we can look at the attributes of the point data using the `@` symbol:

```
head(lnd.stations@data)
```

##	CODE	LEGEND	FILE_NAME	NUMBER	NAME	MICE
## 91	5520	Railway Station	gb_south	17607	Belmont Station	19
## 92	5520	Railway Station	gb_south	17608	Woodmansterne Station	5
## 93	5520	Railway Station	gb_south	17609	Coulsdon South Station	11
## 94	5520	Railway Station	gb_south	17610	Smitham Station	14
## 95	5520	Railway Station	gb_south	17611	Kenley Station	11
## 96	5520	Railway Station	gb_south	17612	Reedham Station	8

In this case we have three potentially interesting variables: “LEGEND”, telling us what the point is, “NAME”, and “MICE”, which represents the number of mice sightings reported by the public at that point (this is a fictional variable). To illustrate the power of the `aggregate` function, let us use it to find the average number of mice spotted in transport points in each London borough, and the standard deviation:

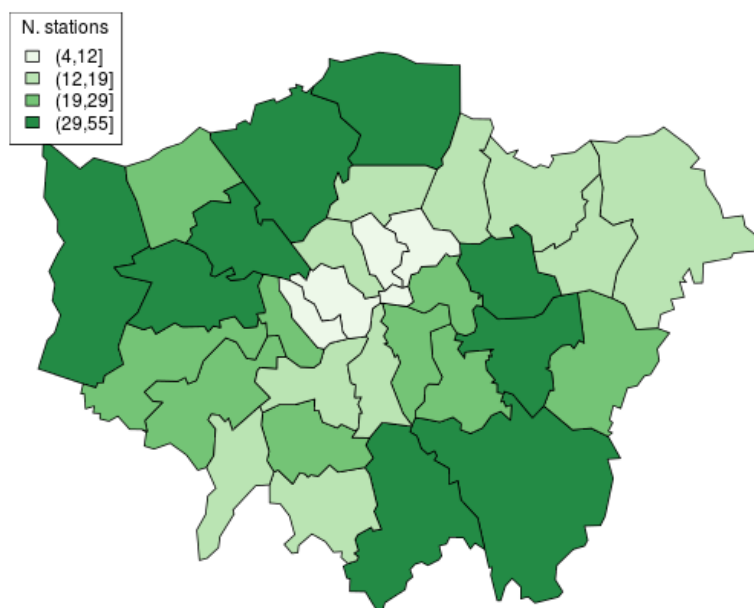


Figure 12: Number of stations in London boroughs

```
lndAvMice <- aggregate(lnd.stations["MICE"], by = lnd, FUN = mean)
summary(lndAvMice)
lndSdMice <- aggregate(lnd.stations["MICE"], by = lnd, FUN = sd)
summary(lndSdMice)
```

Clipping

References

- Bivand, R., & Gebhardt, A. (2000). Implementing functions for spatial statistical analysis using the language. *Journal of Geographical Systems*, 2(3), 307–317.
- Bivand, R. S., Pebesma, E. J., & Rubio, V. G. (2008). *Applied spatial data: analysis with R*. Springer.
- Burrough, P. A. & McDonnell, R. A. (1998). *Principals of Geographic Information Systems* (revised edition). Clarendon Press, Oxford.
- Goodchild, M. F. (2007). Citizens as sensors: the world of volunteered geography. *GeoJournal*, 69(4), 211–221.
- Harris, R. (2012). A Short Introduction to R. social-statistics.org.
- Kabacoff, R. (2011). *R in Action*. Manning Publications Co.
- Krygier, J. Wood, D. 2011. *Making Maps: A Visual Guide to Map Design for GIS* (2nd Ed.). New York: The Guildford Press.
- Longley, P., Goodchild, M. F., Maguire, D. J., & Rhind, D. W. (2005). *Geographic information systems and science*. John Wiley & Sons.
- Monkhouse, F.J. and Wilkinson, H. R. (1973). *Maps and Diagrams Their Compilation and Construction* (3rd Edition, reprinted with revisions). London: Methuen & Co Ltd.
- Ramsey, P., & Dubovsky, D. (2013). Geospatial Software’s Open Future. *GeoInformatics*, 16(4).
- Sherman, G. (2008). *Desktop GIS: Mapping the Planet with Open Source Tools*. Pragmatic Bookshelf.
- Torfs and Brauer (2012). A (very) short Introduction to R. The Comprehensive R Archive Network.
- Venables, W. N., Smith, D. M., & Team, R. D. C. (2013). An introduction to R. The Comprehensive R Archive Network (CRAN). Retrieved from <http://cran.ma.imperial.ac.uk/doc/manuals/r-devel/R-intro.pdf>.
- Wickham, H. (2009). *ggplot2: elegant graphics for data analysis*. Springer.

Wickham, H. (2010). A Layered Grammar of Graphics. American Statistical Association, Institute of Mathematics Statistics and Interface Foundation of North America Journal of Computational and Graphical Statistics. 19, 1: 3-28.

Endnotes

1. R's name originates from the creators of R, Ross Ihaka and Robert Gentleman. R is an open source implementation of the statistical programming language S, so its name is also a play on words that makes implicit reference to this.
2. R is notoriously difficult to search for on major search engines, as it is such a common letter with many other uses beyond the name of a statistical programming language. This should not be a deterrent, as R has a wealth of excellent online resources. To overcome the issue, you can either be more specific with the search term (e.g. "R spatial statistics") or use an R specific search engine such as rseek.org. You can also search of online help *from within R* using the command `RSiteSearch`. E.g. `RSiteSearch("spatial statistics")`. Experiment and see which you prefer!