

Spatial data visualisation with R

Cheshire, James
james.cheshire@ucl.ac.uk

Lovelace, Robin
r.lovelace@leeds.ac.uk

December 24, 2013

Introduction

What is R?

R is a free and open source computer program that runs on all major operating systems. R relies primarily on the *command line* for data input: instead of interacting with the program by moving your mouse around clicking on different parts of the screen, users enter commands via the keyboard. This will seem strange to people accustomed to relying on a graphical user interface (GUI) for most of their computing, e.g. via popular programs such as Microsoft Excel or SPSS, yet the approach has a number of benefits, as highlighted by Gary Sherman (2008, p. 283), developer of the popular Geographical Information System QGIS:

With the advent of “modern” GIS software, most people want to point and click their way through life. That’s good, but there is a tremendous amount of flexibility and power waiting for you with the command line. Many times you can do something on the command line in a fraction of the time you can do it with a GUI.

The joy of this, when you get accustomed to it, is that any command is only ever a few keystrokes away, and the order of the commands sent to R can be stored and repeated in scripts, saving even more time in the long-term (more on this in section ...).

Another important attribute of R, related to its command line interface, is that it is a fully fledged *programming language*. Other GIS programs are written in lower level languages such as C++ which are kept at a safe distance from the users by the GUI. In R, by contrast, the user is ‘close to the metal’ in the sense that what he or she inputs is the same as what R sees when it processes the request. This ‘openness’ can seem raw and daunting to beginners, but it is vital to R’s success. Access to R’s source code and openness about how it works has enabled a veritable army of programmers to improve R over time and add an incredible number of extensions to its base capabilities. Consider for a moment that there are now more than 4000 official packages for R, allowing it to tackle almost any computational or numerical problem one could image, and many more that one could not!

Although writing R source code and creating new packages will not appeal to most R users, it inspires confidence to know that there is a strong and highly skilled community of R developers. If there is a useful spatial function that R cannot currently perform, there is a reasonable chance that someone is working on a solution that will become available at a later date. This constant evolution and improvement is a feature of open source software projects not limited to R, but the range and diversity of extensions is certainly one of its strong points. One area where extension of R’s basic capabilities has been particularly successful is the addition of a wide variety of spatial tools.

The rise of R’s spatial capabilities

!!! Quick history of R’s spatial packages emphasizing current growth and heavy dependence on sp.

!!!Mention exciting and recently added packages.

!!! Not sure if this is worth a complete section- we probably mention many of the things throughout the chapter. I also wonder who really cares!

Why R for spatial data visualisation?

!!! Add a bit about GIS and also the rise of information graphics- arguably GIS has been caught out.

Aside from confusion surrounding its one character name - “what kind of a name is R?” [1] and “how can you possibly find resources for R online?” [2] - R may also seem a strange choice for a chapter on *spatial* data visualisation specifically. “I thought R was just for statistics?” and “Why not use a proper GIS package like ArcGIS?” are valid questions.

The first question arises because R was traditionally conceived - and is still primarily known - as a “statistical programming language” (Bivand and Gebhardt 2000). Although R does have cutting edge statistical capabilities, this definition does not do justice to its power and flexibility. Thus, a more accurate albeit longer definition of R is “an integrated suite of software facilities for data manipulation, calculation and graphical display” (Venables et al. 2013). It is important to consider this wider definition before diving into R: it is a fully fledged programming language meaning that it is highly extensible but also that the same result can often be generated in different ways. This can be confusing.

The second question is based on the premise that all ‘proper’ GISs need to operate in the same way, with primacy allocated to a mapping window and a mouse-driven GUI interface. But when we look back at the history of GIS and its definitions, it becomes clear that R *is* fully fledged GIS, when it is set up correctly. All early GIS programs used a command-line interface; GUIs were only developed later as a way to run commands without needing to remember all the command names (although this is largely overcome by good ‘help’ options and auto-completion). A concise definition of a GIS is “a computerized tool for solving geographic problems” (Longley et al. 2005, p. 16) and R certainly enables this. A more expansive definition of GIS is “a powerful set of tools for collecting, storing, retrieving at will, transforming, and displaying spatial data from the real world for a particular set of purposes” (Burrough and McDonnell, 1998, from Bivand et al. 2013, p. 5); R excels at each of these tasks.

That being said, there are a few major differences between R and conventional GIS programs in terms of spatial data visualisation: R is more suited to creating one-off graphics than exploring spatial data through repeated zooming, panning and spatial sub-setting using custom-drawn polygons, compared with conventional GIS programs. Although interactive maps in R can be created (e.g. using the web interface `shiny`), it is recommended that R is used *in addition to* rather than as a direct replacement of dedicated GIS programs, especially now that there are myriad free options to try (Sherman 2008). An additional point is that while dedicated GIS programs handle spatial data by default and display the results in a single way, there are various options in R that must be decided by the user, for example whether to use R’s base graphics or a dedicated graphics package such as `ggplot2`. On the other hand, the main benefits of R for spatial data visualisation lie in the *reproducibility* of its outputs, a feature that we will be using to great effect in this chapter.

Finally, there continues to be a clear drive towards open and transparent datasets and this chimes well with the open access agenda gaining momentum in academic publishing. R encourages truly transparent and reproducible academic research by enabling anyone with an R installation to re-run the code written to produce the results described in a paper. This process is facilitated by the encouragement of more consistent syntax structures and also developments in the RStudio integrated development environment (IDE) that offers formats such as `.rmd` that are executable R files in addition to the means of sharing and documenting code widely through Git and Subversion systems.

R in the wild

Examples of where R has had an important visual impact.

Might be good to mention New York Times etc here as key users of R.

Chapter overview

The remainder of this chapter will introduce a series of important concepts associated with sound spatial data visualisation. It will provide examples of R code to illustrate each concept and offer inspiration for work with your own data.

R and Spatial Data

Spatial Data in R

In any data analysis project, spatial or otherwise, it is important to have a strong understanding of the dataset before progressing. This section will therefore begin with a description of the input data. We will see how data can be loaded into R and exported to other formats, before going into more detail about the underlying structure of spatial data in R: how it ‘sees’ spatial data is quite unique.

Loading spatial data in R

In most situations, the starting point of spatial analysis tasks is loading in datasets. These may originate from government agencies, remote sensing devices or ‘volunteered geographical information’ (Goodchild 2007). The diversity of geographical data formats is large.

R is able to import a very wide range of spatial data formats thanks to its interface with the Geospatial Data Abstraction Library (GDAL), which is enabled by the package `rgdal`. Below we will load data from two spatial data formats: GPS eXchange (`.gpx`) and an ESRI Shapefile.

`readOGR` is in fact capable of loading dozens more file formats, so the focus is on the *method* rather than the specific formats. The ‘take home message’ is that the `readOGR` function is capable of loading most common spatial file formats, but behaves differently depending on file type. Let’s start with a `.gpx` file, a tracklog recording a bicycle ride from Sheffield to Wakefield which was uploaded Open Street Map [3].

```
# download.file('http://www.openstreetmap.org/trace/1619756/data', destfile
# = 'data/gps-trace.gpx')
library(rgdal) # load the gdal package
shf2lds <- readOGR(dsn = "data/gps-trace.gpx", layer = "tracks") # load track
plot(shf2lds)
shf2lds.p <- readOGR(dsn = "data/gps-trace.gpx", layer = "track_points") # load points
points(shf2lds.p[seq(1, 3000, 100), ])
```




Figure 1: Leeds to Sheffield GPS data with latitude and longitude axes

There is a lot going on in the preceding 7 lines of code, including functions that you are unlikely to have encountered before. Let us think about what has happened, line-by-line.

First, we used R to *download* a file from the internet, using the function `download.file`. The two essential arguments of this function are `url` (we could have typed `url =` before the link) and `destfile`, the destination file. As with any function, more optional arguments can be viewed by by typing `?download.file`.

When `rgdal` has successfully loaded, the next task is not to import the file directly, but to find out which *layers* are available to import, with `ogrListLayers`. The output from this command tells us that various layers are available, including `tracks` and `track_points`: try it. These are imported into R's *workspace* using `readOGR`.

Finally, the basic `plot` function is used to visualize the newly imported objects, ensuring they make sense. In the second `plot` function, we take a subset of the object (see section ... for more on this). To see how to add axes, enter `?axis`

As stated in the help documentation (accessed by entering `?readOGR`), the `dsn =` argument is interpreted differently depending on the type of file used. In the above example, the file name was the file name. To load Shapefiles, by contrast, the *folder* containing the data is used:

```
lnd <- readOGR(dsn = "data/", "london_sport")
```

Here, the files reside in a folder entitled `data` which in R's current working directory (remember to check this using `getwd()`). If the files were stored in the working directory, one would use `dsn = "."` instead. Again, it may be wise to plot the data that results, to ensure that it has worked correctly. Now that the data has been loaded into R's own `sp` format, try interrogating and plotting it, using functions such as `summary` and `plot`.

The size of spatial datasets in R

Any data that has been read into R's *workspace*, which constitutes all objects that can be accessed by name and can be listed using the `ls()` function, can be saved in R's own data storage file type, `.RData`. Spatial datasets can get quite large and this can cause problems on computers by consuming all available memory (RAM) or hard disk space. It is wise to understand roughly how large spatial objects are, providing insight into how long certain functions will take to run.

In the absence of prior knowledge, which of the two objects loaded in the previous section would one expect to be larger? One could hypothesize that the London boroughs represented by the object `lnd` would be larger based on its greater spatial extent, but how much larger? The answer in R is found in the function `object.size`:

```
object.size(shf2lds)
```

```
## 103168 bytes
```

```
object.size(lnd)
```

```
## 79168 bytes
```

Surprisingly, the GPS data is larger. To see why, we can find out how many *vertices* (points connected by lines) are contained in each dataset:

```
shf2lds.f <- fortify(shf2lds)
nrow(shf2lds.f)
```

```
## [1] 6085
```

```
lnd.f <- fortify(lnd)
nrow(lnd.f)
```

```
## [1] 1102
```

In the above block of code we performed two functions for each object: 1) *flatten* the dataset so that each vertice is allocated a unique row 2) use `nrow` to count the result. Note the use of the `<-` assignment operator to create new objects with a `.f` ending.

It is clear that the GPS data has almost 6 times the number of vertices as does the London data, explaining its larger size. Yet when plotted, the GPS data does not seem more detailed, implying that some of the vertices in the object are not needed for visualisation at the scale of the objects *bounding box*.

Simplifying geometries

The wastefulness of the GPS data for visualisation (the full dataset may be useful for other types of analysis) raises the question following question: can the object be simplified such that its key features remain while substantially reducing its size? The answer is yes. In the code below, we harness the power of the `rgeos` package and its `gSimplify` function to simplify spatial R objects:

```
library(rgeos)
shf2lds.simple <- gSimplify(shf2lds, tol = 0.001)
(object.size(shf2lds.simple)/object.size(shf2lds))[1]
```

```
## [1] 0.03047
```

```
plot(shf2lds.simple)
plot(shf2lds, col = "red", add = T)
```

In the above block of code, `gSimplify` is given the object `shf2lds` and the `tol` argument of 0.001 (much larger tolerance values may be needed, for data that is *projected*). Next, we divide the size of the simplified object by the original (note the use of the `/` symbol). The output of 0.03... tells us that the new object is only 3% of its original size. We can see how this has happened by again counting the number of vertices. This time we use the `coordinates` and `nrow` functions together:

```
nrow(coordinates(shf2lds.simple)[[1]][[1]])
```

```
## [1] 44
```

The syntax of the double square brackets will seem strange, providing a taster of how R ‘sees’ spatial data (see section x). Do not worry about this for now. Of interest here is that the number of vertices has shrunk, from 6,084 to only 44, without losing much information about the shape of the line. To test this, try plotting the original and simplified tracks on your computer: when visualized using the `plot` function, object `shf2lds.simple` retains the overall shape of the line and is virtually indistinguishable from the original object.

This example is rather contrived because even the larger object `shf2lds` is only 0.107 Mb, negligible compared with the gigabytes of RAM available to modern computers. However, it underlines a wider point: for visualizing *small scale* maps, spatial data *geometries* can often be simplified to reduce processing time and use of computer memory.

Saving and exporting spatial objects

A typical R workflow involves loading the data, processing and finally exporting the data in a new form. `writeOGR`, the logical counterpart of `readOGR` is ideal for this task. Imagine that we want to view the simplified `gpx` data in software that can only read Shapefiles. This is performed using the following command:

```
shf2lds.simple <- SpatialLinesDataFrame(shf2lds.simple, data.frame(row.names = "0",
  a = 1))
writeOGR(shf2lds.simple, layer = "shf2lds", dsn = "data/", driver = "ESRI Shapefile")
```

In the above code, the object was first converted into a spatial dataframe class required by the `writeOGR` command, before being exported as a shapefile entitled `shf2lds`. Unlike with `readOGR`, the driver must be specified, in this case with “ESRI Shapefile” [4]. The simplified GPS data is now available to other GIS programs for further analysis. Alternatively, `save(shf2lds.simple, file = "data/shf2lds.RData")` will save the object in R’s own spatial data format, which is described in the next section.

The structure of spatial data in R

Spatial datasets in R are saved in their own format, defined as `Spatial...` classes within the `sp` package. For this reason, `sp` is the basic spatial package in R, upon which the others depend. Spatial classes range from the basic `Spatial` class to the complex, `SpatialPolygonsDataFrame`: the `Spatial` class contains only two required `slots`[5]:

```
getSlots("Spatial")
```

```
##          bbox proj4string
## "matrix"      "CRS"
```

This tells us that `Spatial` objects must contain a bounding box (`bbox`) and a CRS. Further details on these can be found by typing `?bbox` and `?proj4string`. All other spatial classes in R build on this foundation of a bounding box and a projection system (which is set automatically to `NA` if it is not known). However, more complex classes contain more slots, some of which are lists which contain additional lists. To find out the slots of `shf2lds.simple`, for example, we would first ascertain its class and then use the `getSlots` command:

```
class(shf2lds.simple) # identify the object's class
```

```
## [1] "SpatialLinesDataFrame"
## attr(,"package")
## [1] "sp"
```

```
getSlots("SpatialLinesDataFrame") # find the associated slots
```

```
##          data          lines          bbox proj4string
## "data.frame"      "list"      "matrix"      "CRS"
```

The same principles apply to all spatial classes including `Spatial*` `Points`, `Polygons`, `Grids` and `Pixels` as well as associated `*DataFrame` classes. For more information on this, see the `sp` documentation: `?Spatial`.

Manipulating spatial data

Coordinate reference systems

As mentioned in the previous section, all `Spatial` objects in R are allocated a coordinate reference system (CRS). The CRS of any spatial object can be found using the command `proj4string`. In some cases the CRS is not known: in this case the result will simply be `NA`. To discover the CRS of the `lnd` object for example, type the following:

```
proj4string(lnd)
```

```
## [1] "+proj=tmerc +lat_0=49 +lon_0=-2 +k=0.9996012717 +x_0=400000 +y_0=-100000 +ellps=airy "
```

The output may seem cryptic but is in fact highly informative: `lnd` has *projected* coordinates, based on the *Transverse Mercator* system (hence `"+proj=tmerc"` in the output) and its origin is at latitude 49N, -2E. This point is

If we *know* that the CRS is incorrectly specified, it can be re-set. In this case, for example we know that `lnd` actually has a CRS OSGB1936. Knowing also that the code for this is 27700, it can be updated as follows:

```
proj4string(lnd) <- CRS("+init=epsg:27700")
```

The CRS has now been updated - note that the key details are all the same as before. Note: this method **should never** be used as an attempt to *reproject* data from one CRS to another.

Reprojecting data

Transforming the coordinates of spatial data from one CRS to another (reprojection) is a common task in GIS. This is because data from national sources are generally provided in *projected* coordinates (the location on the cartesian coordinates of a map) whereas data from GPSs and the internet are generally provided in *geographic* coordinates, with latitude and longitude measured in degrees to locate points on the surface of the globe.

Reprojecting data in R is quite simple: all you need is a spatial object with a known CRS and knowledge of the CRS you wish to transform it to. To illustrate why that is necessary, try to plot the objects `lnd` and `shf2lnd.simple` on the same map:

```
combined <- rbind(fortify(shf2lds.simple)[, 1:2], fortify(lnd)[, 1:2])
plot(combined)
```

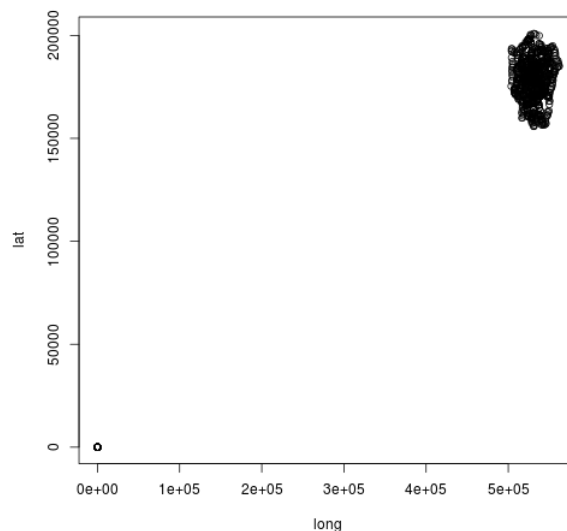


Figure 2: Plot of spatial objects with different CRS

In the above code we first extracted the coordinates of each point using `fortify` and then plotted them using `plot`. The image shows why reprojection is necessary: the `.gpx` data are on a totally different scale than the shapefile of London. Hence the tiny dot at the bottom right of the graph. We will now reproject the data, allowing `lnd` and `shf2lds.simple` to be usefully plotted on the same graphic:

```
lnd.wgs84 <- spTransform(lnd, CRSobj = CRS("+init=epsg:4326"))
```

The above code created a new object, `lnd.wgs84`, that contains the same geometries as the original but in a new CRS using the `spTransform` function. The CRS argument was set to `"+init=epsg:4326"`, which represents the WGS84 CRS via an EPSG code [6]. Now `lnd` has been reprojected we can plot it next to the GPS data:

```
combined <- rbind(fortify(shf2lds.simple)[, 1:2], fortify(lnd.wgs84)[, 1:2])
plot(combined)
```

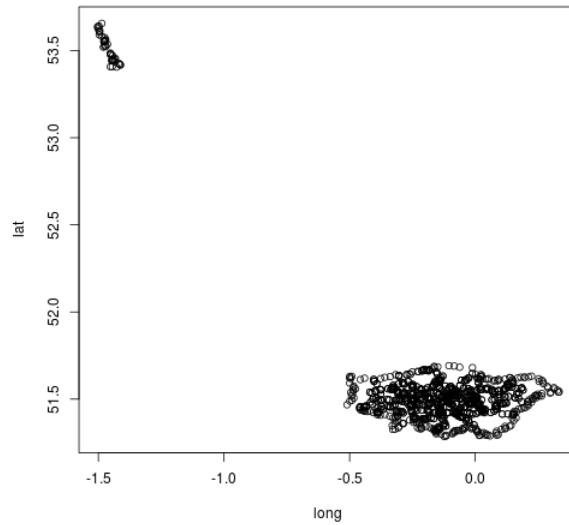


Figure 3: Plot of spatial objects sharing the same CRS

Although the plot of the reprojected data is squashed because the axis scales are not fixed and distorted (*geographic* coordinates such as WGS84 should not usually be used for plotting), at least the relative position and shape of both objects can now be seen. The presence of the dotted line in the top left of the plot confirms our assumption that the GPS data is from around Sheffield, which is northwest of London.

Attribute joins

Because boroughs are official administrative zones, there is much data available at this level that we can link to the polygons in the `lnd` object. We will use the example of crime data to illustrate this data availability, which is stored in the `data` folder available from this project's github page.

```
load("data/crimeAg.Rdata") # load the crime dataset from an R dataset
```

After the dataset has been explored (e.g. using the `summary` and `head` functions) to ensure compatibility, it can be joined to `lnd`. We will use the `join` function in the `plyr` package but the `merge` function could equally be used (remember to type `library(plyr)` if needed).

`join` requires all joining variables to have the same name, but this work has already been done [7]. Once this preparation has been done, the join function is actually very simple:

```
lnd@data <- join(lnd@data, crimeAg)
```

Take a look at the `lnd@data` object. You should see new variables added, meaning the attribute join was successful.

Spatial joins

A spatial join, like attribute joins, is used to transfer information from one dataset to another. There is a clearly defined direction to spatial joins, with the *target layer* receiving information from another spatial layer based on the proximity of elements from both layers to each other. There are three broad types of spatial join: one-to-one, many-to-one and one-to-many. We will focus only the former two as the third type is rarely used.

One-to-one spatial joins

One-to-one spatial joins are by far the easiest to understand and compute because they simply involve the transfer of attributes in one layer to another, based on location. A one-to-one join is depicted in figure x below, and can be performed using the same technique as described in the section on spatial aggregation.



Figure 4: Illustration of a one-to-one spatial join

Many-to-one spatial joins

Many-to-one spatial joins involve taking a spatial layer with many elements and allocating the attributes associated with these elements to relatively few elements in the target spatial layer. A common type of many-to-one spatial join is the allocation of data collected at many point sources unevenly scattered over space to polygons representing administrative boundaries, as represented in Fig. x.

```
lnd.stations <- readOGR("data/", "lnd-stns", p4s = "+init=epsg:27700")

## OGR data source with driver: ESRI Shapefile
## Source: "data/", layer: "lnd-stns"
## with 2532 features and 6 fields
## Feature type: wkbPoint with 2 dimensions

plot(lnd)
plot(lnd.stations[round(runif(500, 1, nrow(lnd.stations))), ], add = T)
```

The above code reads in a `SpatialPointsDataFrame` consisting of 2532 transport nodes in and surrounding London and then plots a random sample of 500 of these over the previously loaded borough level administrative boundaries. The reason for piloting a sample of the points rather than all of them is that the boundary data becomes difficult to see if all of the points are plotted. It is also useful to see and practice sampling techniques in practice; try to plot only the first 500 points, rather than a random selection, and describe the difference.

The most obvious issue with the point data from the perspective of a spatial join with the borough data is that many of the points in the dataset are in fact located outside the region of interest. Thus, the first stage in the analysis is to filter the point data such that only those that lie within London's administrative zones are selected. This in itself is a kind of spatial join, and can be accomplished with the following code.




Figure 5: Input data for a spatial join

```
proj4string(lnd) <- proj4string(lnd.stations)
lnd.stations <- lnd.stations[lnd, ] # select only points within lnd
plot(lnd.stations) # check the result
```




Figure 6: A spatial subset of the points

The station points now clearly follow the form of the `lnd` shape, indicating that the procedure worked. Let's review the code that allowed this to happen: the first line ensured that the CRS associated with each layer is *exactly* the same: this step should not be required in most cases, but it is worth knowing about. Of course, if the coordinate systems are *actually* different in each layer, the function `spTransform` will be needed to make them compatible. This procedure is discussed in section !!! In this case, only the name was slightly different hence direct alteration of the CRS name via the function `proj4string`.

The second line of code is where the magic happens and the brilliance of R's `sp` package becomes clear: all that was needed was to place another spatial object in the row index of the points (`[lnd,]`) and R automatically understood that a subset based on location should be produced. This line of code is an example of R's 'terseness' - only a single line of code is needed to perform what is in fact quite a complex operation.

Spatial aggregation

Now that only stations which *intersect* with the `lnd` polygon have been selected, the next stage is to extract information about the points within each zone. This many-to-one spatial join is also known as *spatial aggregation*. To do this there are a couple of approaches: one using the `sp` package and the other using `rgeos` (see Bivand et al. 2013, 5.3).

As with the *spatial subset* method described above, the developers of R have been very clever in their implementation of spatial aggregations methods. To minimise typing and ensure consistency with R's base functions, `sp` extends the capabilities of the `aggregate` function to automatically detect whether the user is asking for a spatial or a non-spatial aggregation (they are, in essence, the same thing - we recommend learning about the non-spatial use of `aggregate` in R for comparison).

Continuing with the example of station points in London polygons, let us use the spatial extension of `aggregate` to count how many points are in each borough:

```
lndStC <- aggregate(lnd.stations, by = lnd, FUN = length)
summary(lndStC)
plot(lndStC)
```

As with the spatial subset function, the above code is extremely terse. The `aggregate` function here does three things: 1) identifies which stations are in which London borough; 2) uses this information to perform a function on the output, in this case `length`, which simply means "count" in this context; and 3) creates a new spatial object equivalent to `lnd` but with updated attribute data to reflect the results of the spatial aggregation. The results, with a legend and colours added, are presented in Fig !!! below.




Figure 7: Number of stations in London boroughs

As with any spatial attribute data stored as an `sp` object, we can look at the attributes of the point data using the `@` symbol:

```
head(lnd.stations@data, n = 2)
```

##	CODE	LEGEND	FILE_NAME	NUMBER	NAME	MICE
## 91	5520	Railway Station	gb_south	17607	Belmont Station	19
## 92	5520	Railway Station	gb_south	17608	Woodmansterne Station	5

In this case we have three potentially interesting variables: “LEGEND”, telling us what the point is, “NAME”, and “MICE”, which represents the number of mice sightings reported by the public at that point (this is a fictional variable). To illustrate the power of the `aggregate` function, let us use it to find the average number of mice spotted in transport points in each London borough, and the standard deviation:

```
lndAvMice <- aggregate(lnd.stations["MICE"], by = lnd, FUN = mean)
lndSdMice <- aggregate(lnd.stations["MICE"], by = lnd, FUN = sd)
```

In the above code, `aggregate` was used to create entirely new spatial objects that are exactly the same as `lnd`, except with new attribute data. To add the mean mice count to the original object, the following code can be used:

```
lnd$av.mice <- lndAvMice$MICE
```

The above code creates a new variable in the `lnd@data` object entitled “av.mice” and populates it with desired values. Thus `Spatial` objects can behave in the same way as `data.frames` when referring to attribute variables.

Summary

To summarise this section, we have taken a look inside R’s representation of spatial data, learned how to manipulate these datasets in terms of CRS transformations and attribute data and finally explored spatial joins and aggregation. Only *after* the datasets are well understood and saved in a suitable form should we move on to visualisation.

Fundamentals of Spatial Data Visualisation

Good maps depend on sound analysis and data preparation and can have an enormous impact on the understanding and communication of results. It has never been easier to produce a map. The underlying data required are available in unprecedented volumes and the technological capabilities of transforming them into compelling maps and graphics are increasingly sophisticated and straightforward to use. Data and software, however, only offer the starting points of good spatial data visualisation since they need to be refined and calibrated by the researchers seeking to communicate their findings. In this section we will run through the features of a good map. We will then seek to emulate them with R in Section XX. It is worth noting that not all good maps and graphics contain all the features below – they should simply be seen as suggestions rather than firm principles.

Effective map making is hard process – as Krygier and Wood (XXX) put it “there is a lot to see, think about, and do” (p6). It often comes at the end of a period of intense data analysis and perhaps when the priority is to get a paper finished or results published and can therefore be rushed as a result. The beauty of R (and other scripting languages) is the ability to save code and simply re-run it with different data. Colours, map adornments and other parameters can therefore be quickly applied so it is well worth creating a template script that adheres to best practice.

We have selected `ggplot2` as our package of choice for the bulk of our maps and spatial data visualisations because it has a number of these elements at its core. The “gg” in its slightly odd name stands for “Grammar of Graphics”, which is a set of rules developed by Leland Wilkinson (2005) in a book of the same name. Grammar in the context of graphics works in much the same way as it does in language- it provides a structure. The structure is informed by both human perception and also mathematics to ensure that the resulting visualisations are both technically sound and comprehensible. Through creating `ggplot2`, Hadley Wickham, implemented these rules as well as developing ways in which plots can be built up in layers (see Wickham, 2010). This layering component is especially useful in the context of spatial data since it is conceptually the same as map layers in Geographical Information Systems (GIS).

!!!!Maps with `ggplot2` !!!!Adding base maps with `ggmap`

First load the libraries required for this section:

```
library(rgdal)

## Loading required package: sp
## rgdal: version: 0.8-10, (SVN revision 478)
## Geospatial Data Abstraction Library extensions to R successfully loaded
## Loaded GDAL runtime: GDAL 1.10.0, released 2013/04/24
## Path to GDAL shared files: /usr/share/gdal/1.10
## Loaded PROJ.4 runtime: Rel. 4.8.0, 6 March 2012, [PJ_VERSION: 480]
## Path to PROJ.4 shared files: (autodetected)

library(ggplot2)
library(gridExtra)

## Loading required package: grid
```

You will also need create a folder and then set it as your working directory. Assuming your name is **Uname**, and the folder is saved as **sdvwR** in the Desktop in Windows, use the following.

```
setwd("c:/Users/Uname/Desktop/sdvwR")
```

For this section we are going to use a map of the world to demonstrate some of the cartographic principles discussed. A world map is available from the Natural Earth website. The code below will download this and save it to your working directory. It is commented out because the data may already be on your system. Uncomment each new line (by deleting the # symbol) if you need to download and extract the data.

```
download.file(url = "http://www.naturalearthdata.com/http://www.naturalearthdata.com/download/110m/cultural.
  "ne_110m_admin_0_countries.zip", "auto") # download file
unzip("ne_110m_admin_0_countries.zip", exdir = "data/") # unzip to data folder
file.remove("ne_110m_admin_0_countries.zip") # remove zip file
```

Once downloaded we can then load the data into the R console. We have just downloaded a shapefile, which as Section XX explains, is not handled as a “standard” data object in R.

```
wrld <- readOGR("data/", "ne_110m_admin_0_countries")

## OGR data source with driver: ESRI Shapefile
## Source: "data/", layer: "ne_110m_admin_0_countries"
## with 177 features and 63 fields
## Feature type: wkbPolygon with 2 dimensions

plot(wrld)
```

To see the first ten rows of attribute information associated with each of the country boundaries type the following:

```
head(wrld@data)[, 1:5]
```

```
##  scalerank      featurecla labelrank      sovereignty sov_a3
##  0          1 Admin-0 country      3      Afghanistan  AFG
##  1          1 Admin-0 country      3          Angola  AGO
##  2          1 Admin-0 country      6          Albania  ALB
##  3          1 Admin-0 country      4 United Arab Emirates ARE
##  4          1 Admin-0 country      2          Argentina ARG
##  5          1 Admin-0 country      6          Armenia  ARM
```



Figure 8: unnamed-chunk-4

You can see there are a lot of columns associated with this file. Although we will keep all of the them, we are only really interested in the population estimate (“pop_est”) field. Before progressing it is worth reprojecting the data in order that the population data can be seen better. The coordinate reference system of the wrld shapefile is currently WGS84. This the common latitude and longitude format that all spatial software packages understand. From a cartographic perspective the standard plots of this projection, of the kind produced above, are not suitable since they distort the shapes of those countries further from the equator. Instead the Robinson projection provides a good compromise between areal distortion and shape preservation. We therefore project it as follows.

```
wrld.rob <- spTransform(wrld, CRS("+proj=robin"))
plot(wrld.rob)
```



Figure 9: unnamed-chunk-6

“+proj=robin” refers to the Robinson projection. You will have spotted from the plot that the countries in the world map are much better proportioned.

We now need to “fortify” this spatial data to convert it into a format that ggplot2 understands, we also use “merge” to re-attach the attribute data that is lost in the fortify operation. !!! explain fortify

```
wrld.rob.f <- fortify(wrld.rob, region = "sov_a3")

## Loading required package: rgeos
## rgeos version: 0.2-19, (SVN revision 394)
## GEOS runtime version: 3.3.8-CAPI-1.7.8
## Polygon checking: TRUE

wrld.pop.f <- merge(wrld.rob.f, wrld.rob@data, by.x = "id", by.y = "sov_a3")

# continuous colour ramp

map <- ggplot(wrld.pop.f, aes(long, lat, group = group, fill = pop_est)) + geom_polygon() +
  coord_equal() + labs(x = "Longitude", y = "Latitude", fill = "World Population") +
  ggtitle("World Population")

# better colours with more breaks- to finish

map + scale_fill_continuous(breaks = )

## Error: argument is missing, with no default

# categorical variables
```

Conforming to colour conventions

Colour has an enormous impact on how people will perceive your graphic. “Readers” of a map come to it with a range of pre-conceptions about how the world looks. If the map’s purpose is to clearly communicate data then it is often advisable to conform to conventions so as not to disorientate readers to ensure they can focus on the key messages contained in the data. A good example of this is the use of blue for bodies of water and green for landmass. The code example below generates two plots with our wrld.pop.f object. The first colours the land blue and the sea (in this case the background to the map) green and the second is more conventional. We use the “grid.arrange” function from the “gridExtra” package to display the maps side by side.

```
map2 <- ggplot(wrld.pop.f, aes(long, lat, group = group)) + coord_equal()

blue <- map2 + geom_polygon(fill = "light blue") + theme(panel.background = element_rect(fill = "dark green"))

green <- map2 + geom_polygon(fill = "dark green") + theme(panel.background = element_rect(fill = "light blue"))

grid.arrange(blue, green, ncol = 2)
```

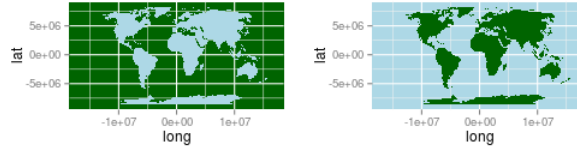


Figure 10: unnamed-chunk-9

Experimenting with line colour and line widths

In addition to conforming to colour conventions, line colour and width offer important parameters, which are often overlooked tools for increasing the legibility of a graphic. As the code below demonstrates, it is possible to adjust line colour through using the “colour” parameter and the line width using the “lwd” parameter. The impact of different line widths will vary depending on your screen size and resolution. If you save the plot to pdf (or an image) then the size at which you do this will also affect the line widths.

```
map3 <- map2 + theme(panel.background = element_rect(fill = "light blue"))

yellow <- map3 + geom_polygon(fill = "dark green", colour = "yellow")

black <- map3 + geom_polygon(fill = "dark green", colour = "black")

thin <- map3 + geom_polygon(fill = "dark green", colour = "black", lwd = 0.1)

thick <- map3 + geom_polygon(fill = "dark green", colour = "black", lwd = 1.5)

grid.arrange(yellow, black, thick, thin, ncol = 2)
```

There are other parameters such as layer transparency that can be applied to all aspects of the plot - both points, lines and polygons - that we will reference in later examples in this chapter.

Map Adornments and Annotations

Map adornments and annotations are essential to orientate the viewer and provide context; they include graticules, north arrows, scale bars and data attribution. Not all are required on a single map, indeed it is often best that they are used sparingly to avoid unnecessary clutter (Monkhouse and Wilkinson, 1971). Unfortunately it is not always as straightforward to add these in R, and perhaps less so using the ggplot2 paradigm, when compared to a conventional GIS. Here we will outline the ways in which annotations can be added.

!!!! In the maps created so far, we have defined the *aesthetics* of the map in the foundation function `ggplot`. The result of this is that all subsequent layers are expected to have the same variables and essentially contain data with

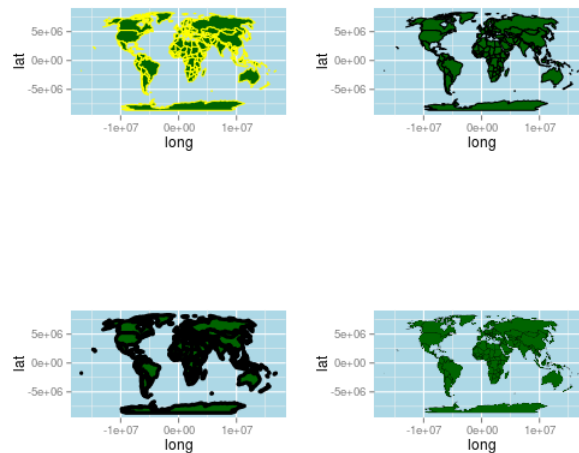


Figure 11: unnamed-chunk-10

the same dimensions as original dataset. But what if we want to add a new layer from a completely different dataset. To do this, we must not add any arguments to the `ggplot` function, only adding data sources one layer at a time:

North arrow

In the maps created so far, we have defined the *aesthetics* of the map in the foundation function `ggplot`. The result of this is that all subsequent layers are expected to have the same variables and essentially contain data with the same dimensions as original dataset. But what if we want to add a new layer from a completely different dataset, e.g. to add an arrow? To do this, we must not add any arguments to the `ggplot` function, only adding data sources one layer at a time:

Here we create an empty plot, meaning that each new layer must be given its own dataset. While more code is needed in this example, it enables much greater flexibility with regards to what can be included in new layer contents. Another possibility is to use the `segment` geom to add a pre-made arrow: !!! This needs to sorted - doesn't compile Robin

```
ggplot() + geom_polygon(data = wrld.pop.f, aes(long, lat, group = group, fill = pop_est)) +
  geom_line(aes(x = c(-160, -160), y = c(0, 25)), arrow = arrow())
```

Scale bar

!!! use `geom_segment` with `geom_text`

There is an almost infinite number of different combinations of the above parameters so take inspiration from maps and graphics you have seen and liked. The process is an iterative one, it will take multiple attempts to get right. Show your map to friends and colleagues- all will have an opinion but don't be afraid to stand by the decisions you have taken.

!!!Consistency- across papers.

`ggmap` is a package that uses the `ggplot2` syntax as a template to create maps with image tiles from the likes of Google and OpenStreetMap:

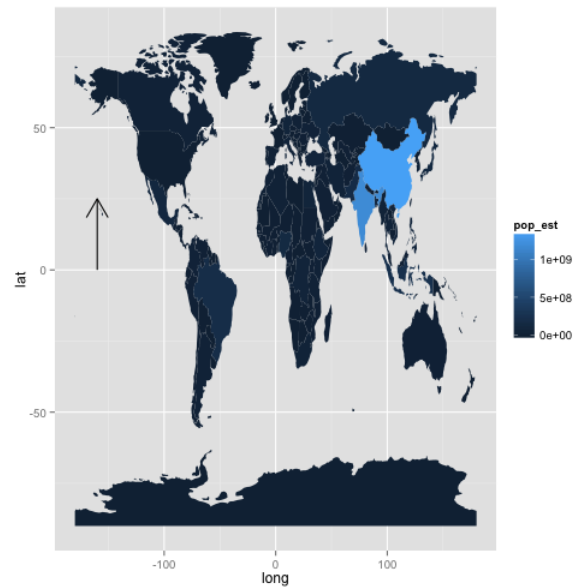


Figure 12: World map with a North arrow

Adding Basemaps To Your Plots

```
library(ggmap)
```

!!! adapt to the lnd.stations object.

The sport object is in British National Grid but the ggmap image tiles are in WGS84. We therefore need to use the sport.wgs84 object created in the reprojection operation earlier.

The first job is to calculate the bounding box (bb for short) of the sport.wgs84 object to identify the geographic extent of the image tiles that we need.

```
b <- bbox(sport.wgs84)
```

```
## Error: error in evaluating the argument 'obj' in selecting a method for function 'bbox': Error: object
```

```
b[1, ] <- (b[1, ] - mean(b[1, ])) * 1.05 + mean(b[1, ])
```

```
## Error: object 'b' not found
```

```
b[2, ] <- (b[2, ] - mean(b[2, ])) * 1.05 + mean(b[2, ])
```

```
## Error: object 'b' not found
```

```
# scale longitude and latitude (increase bb by 5% for plot) replace 1.05  
# with 1.xx for an xx% increase in the plot size
```

This is then fed into the `get_map` function as the location parameter. The syntax below contains 2 functions. `ggmap` is required to produce the plot and provides the base map data.

```
lnd.b1 <- ggmap(get_map(location = b))
```

```
## Error: object 'b' not found
```

In much the same way as we did above we can then layer the plot with different geoms.

First fortify the sport.wgs84 object and then merge with the required attribute data (we already did this step to create the sport.f object).

```
sport.wgs84.f <- fortify(sport.wgs84, region = "ons_label")
```

```
## Error: object 'sport.wgs84' not found
```

```
sport.wgs84.f <- merge(sport.wgs84.f, sport.wgs84@data, by.x = "id", by.y = "ons_label")
```

```
## Error: object 'sport.wgs84.f' not found
```

We can now overlay this on our base map.

```
lnd.b1 + geom_polygon(data = sport.wgs84.f, aes(x = long, y = lat, group = group,
  fill = Partic_Per), alpha = 0.5)
```

The code above contains a lot of parameters. Use the ggplot2 help pages to find out what they are. The resulting map looks okay, but it would be improved with a simpler base map in black and white. A design firm called stamen provide the tiles we need and they can be brought into the plot with the `get_map` function:

```
lnd.b2 <- ggmap(get_map(location = b, source = "stamen", maptype = "toner",
  crop = T))
```

```
## Error: object 'b' not found
```

We can then produce the plot as before.

```
lnd.b2 + geom_polygon(data = sport.wgs84.f, aes(x = long, y = lat, group = group,
  fill = Partic_Per), alpha = 0.5)
```

Finally, if we want to increase the detail of the base map, `get_map` has a `zoom` parameter.

```
lnd.b3 <- ggmap(get_map(location = b, source = "stamen", maptype = "toner",
  crop = T, zoom = 11))
```

```
## Error: object 'b' not found
```

```
lnd.b3 + geom_polygon(data = sport.wgs84.f, aes(x = long, y = lat, group = group,
  fill = Partic_Per), alpha = 0.5)
```

```
## Error: object 'lnd.b3' not found
```

A Final Example

Here we present a final example that draws upon the many advanced concepts discussed in this chapter to produce a map of 18th Century Shipping flows.

```
library(ggplot2)
library(png)
wrld <- readOGR(".", "ne_110m_admin_0_countries")

## Error: Cannot open file

btitle <- readPNG("figure/brit_titles.png")
compass <- readPNG("figure/windrose.png")
earth <- readPNG("figure/earth_raster.png")
bdata <- read.csv("data/british_shipping_example.csv")

xquiet <- scale_x_continuous("", breaks = NULL)
yquiet <- scale_y_continuous("", breaks = NULL)
quiet <- list(xquiet, yquiet)

wrld.pop.f <- fortify(wrld, region = "sov_a3")

## Error: object 'wrld' not found

base <- ggplot(wrld.pop.f, aes(x = long, y = lat))

## Error: object 'wrld.pop.f' not found

route <- c(geom_path(aes(long, lat, group = paste(bdata$trp, bdata$group.regroup,
  sep = ".")), colour = "#0F3B5F", size = 0.2, data = bdata, alpha = 0.5,
  lineend = "round"))

wrld <- c(geom_polygon(aes(group = group), size = 0.1, colour = "black", fill = "#D6BF86",
  data = wrld.pop.f, alpha = 1))

## Error: object 'wrld.pop.f' not found

base + route + wrld + theme(panel.background = element_rect(fill = "#BAC4B9",
  colour = "black")) + annotation_raster(btitle, xmin = 30, xmax = 140, ymin = 51,
  ymax = 87) + annotation_raster(compass, xmin = 65, xmax = 105, ymin = 25,
  ymax = 65) + coord_equal() + quiet

## Error: object 'base' not found

base + annotation_raster(earth, xmin = -180, xmax = 180, ymin = -90, ymax = 90) +
  route + theme(panel.background = element_rect(fill = "#BAC4B9", colour = "black")) +
  annotation_raster(btitle, xmin = 30, xmax = 140, ymin = 51, ymax = 87) +
  annotation_raster(compass, xmin = 65, xmax = 105, ymin = 25, ymax = 65) +
  coord_equal() + quiet

## Error: object 'base' not found

!!!change the colour of the routes to make them stand out more.
```

Recap and Conclusions

References

- Bivand, R., & Gebhardt, A. (2000). Implementing functions for spatial statistical analysis using the language. *Journal of Geographical Systems*, 2(3), 307–317.
- Bivand, R. S., Pebesma, E. J., & Rubio, V. G. (2008). *Applied spatial data: analysis with R*. Springer.
- Burrough, P. A. & McDonnell, R. A. (1998). *Principals of Geographic Information Systems* (revised edition). Clarendon Press, Oxford.
- Goodchild, M. F. (2007). Citizens as sensors: the world of volunteered geography. *GeoJournal*, 69(4), 211–221.
- Harris, R. (2012). A Short Introduction to R. social-statistics.org.
- Kabacoff, R. (2011). *R in Action*. Manning Publications Co.
- Krygier, J. Wood, D. 2011. *Making Maps: A Visual Guide to Map Design for GIS* (2nd Ed.). New York: The Guildford Press.
- Longley, P., Goodchild, M. F., Maguire, D. J., & Rhind, D. W. (2005). *Geographic information systems and science*. John Wiley & Sons.
- Monkhouse, F.J. and Wilkinson, H. R. 1973. *Maps and Diagrams Their Compilation and Construction* (3rd Edition, reprinted with revisions). London: Methuen & Co Ltd.
- Ramsey, P., & Dubovsky, D. (2013). Geospatial Software’s Open Future. *GeoInformatics*, 16(4).
- Sherman, G. (2008). *Desktop GIS: Mapping the Planet with Open Source Tools*. Pragmatic Bookshelf.
- Torfs and Brauer (2012). A (very) short Introduction to R. The Comprehensive R Archive Network.
- Venables, W. N., Smith, D. M., & Team, R. D. C. (2013). An introduction to R. The Comprehensive R Archive Network (CRAN). Retrieved from <http://cran.ma.imperial.ac.uk/doc/manuals/r-devel/R-intro.pdf>.
- Wickham, H. (2009). *ggplot2: elegant graphics for data analysis*. Springer.
- Wickham, H. 2010. A Layered Grammar of Graphics. American Statistical Association, Institute of Mathematics Statistics and Interface Foundation of North America *Journal of Computational and Graphical Statistics*. 19, 1: 3-28

Endnotes

1. R’s name originates from the creators of R, Ross Ihaka and Robert Gentleman. R is an open source implementation of the statistical programming language S, so its name is also a play on words that makes implicit reference to this.
2. R is notoriously difficult to search for on major search engines, as it is such a common letter with many other uses beyond the name of a statistical programming language. This should not be a deterrent, as R has a wealth of excellent online resources. To overcome the issue, you can either be more specific with the search term (e.g. “R spatial statistics”) or use an R specific search engine such as rseek.org. You can also search of online help *from within R* using the command `RSiteSearch`. E.g. `RSiteSearch("spatial statistics")`. Experiment and see which you prefer!
3. For more information about this ride, please see robinlovelace.net.
4. A complete list of drivers for importing and exporting spatial data can be displayed by typing `getGDALDriverNames()`.
5. Slots are elements found ‘inside’ classes of the [S4 object system](#). While the sub-elements of S3 objects such as `data.frame` are referred to using the `$` symbol, the slots of S4 objects are identified using `@`. Thus, the variable `x` of dataframe `df` can be referred to with `df$x`. In the same way, the data associated with a polygon layer such as `lnd` can be accessed with `lnd@data`. Note that `lnd@data` is itself a dataframe, so can be further specified, e.g. with `lnd@data$name`. For more on spatial data classes, see Bivand et al. (2013).

6. EPSG stands for “European Petroleum Survey Group”, but this is not really worth knowing as the organisation is now defunct (www.epsg.org/). The important thing is that EPSG codes provide a unified way to refer to a wide range of coordinate systems, as each CRS has its own epsg code. These can be found at the website spatialreference.org. To see how this website can be useful, try searching for “osgb”, for example to find the epsg code for the British National Grid.
7. To see how the `crimeAg` dataset was created, please refer to the “Creating-maps-in-R” tutorial (Cheshire and Lovelace, 2014) hosted on [GitHub](#). The file “[intro-spatial-rl.pdf](#)” contains this information, in the section on “Downloading additional data”.

```
source("chapter.R") # convert chapter to tex
```