Spatial data visualisation with R

Cheshire, James
james.cheshire@ucl.ac.uk

Lovelace, Robin
r.lovelace@leeds.ac.uk

January 3, 2014

Introduction

What is R?

R is a free and open source computer program that runs on all major operating systems. R relies primarily on the *command line* for data input: instead of interacting with the program by moving your mouse around clicking on different parts of the screen, users enter commands via the keyboard. This will seem to strange to people accustomed to relying on a graphical user interface (GUI) for most of their computing, yet the approach has a number of benefits, as highlighted by Gary Sherman (2008, p. 283), developer of the popular Geographical Informations System QGIS:

> With the advent of "modern" GIS software, most people want to point and click their way through life. That's good, but there is a tremendous amount of flexibility and power waiting for you with the command line. Many times you can do something on the command line in a fraction of the time you can do it with a GUI.

The joy of this, when you get accustomed to it, is that any command is only ever a few keystrokes away, and the order of the commands sent to R can be stored and repeated in scripts, saving even more time in the long-term (more on this in section . . .).

Another important attribute of R, related to its command line interface, is that it is a fully fledged *programming language*. Other GIS programs are written in lower level languages such as C++ which are kept at a safe distance from the users by the GUI. In R, by contrast, the user is 'close to the metal' in the sense that what he or she inputs is the same as what R sees when it processes the request. This 'openness' can seem raw and daunting to beginners, but it is vital to R's success. Access to R's source code and openness about how it works has enabled a veritable army of programmers to improve R over time and add an incredible number of extensions to its capabilities. There are now more than 4000 official packages for R, allowing it to tackle almost any computational or numerical problem one could imagine.

Although writing R source code and creating new packages will not appeal to most R users, it inspires confidence to know that there is a strong and highly skilled community of R developers. If there is a useful spatial function that R cannot currently perform, there is a reasonable chance that someone is working on a solution that will become available at a later date. This constant evolution and improvement is a feature of open source software projects not limited to R, but the range and diversity of extensions is certainly one of its strong points. One area where extension of R's basic capabilities has been particularly successful is the addition of a wide variety of spatial tools.

Why R for spatial data visualisation?

Aside from confusion surrounding its one character name [1] and uncertainty about how to search for help [2], R may also seem a strange choice for a chapter on *spatial* data visualisation specifically. "I thought R was just for statistics?" and "Why not use a proper GIS package like ArcGIS?" are valid questions.

R was conceived - and is still primarily known - for its capabilities as a "statistical programming language" (Bivand and Gebhardt 2000). Statistical analysis functions remain core to the package but there is a broadening

of functionality to reflect a growing user base across disciplines. R has become "an integrated suite of software facilities for data manipulation, calculation and graphical display" (Venables et al. 2013). Spatial data analysis and visualisation is an important growth area within this increased functionality to the extent that R can almost entirely replace major GIS packages for a whole host of spatial analysis workflows. That said, it will never be a complete replacement for those seeking a user interface that enables panning and zooming, or for those seeking to manually digitise spatial data. We therefore suggest using R *in addition to* GIS packages where required. In fact, R's spatial capabilities have already been integrated in conventional GIS packages ArcGIS (via its R toolbox and the Geospatial Modelling Environment) and QGIS (via the *Processing* framework).

R and conventional GIS programs

There are a few major differences between R and conventional GIS programs in terms of spatial data visualisation: R is more suited to creating one-off graphics than exploring spatial data through repeated zooming, panning and spatial sub-setting using custom-drawn polygons, compared with conventional GIS programs. Although interactive maps in R can be created (e.g. using the web interface `shiny`), it is recommended that R is used *in addition to* rather than as a direct replacement of dedicated GIS programs, especially now that there are myriad free options to try (Sherman 2008). An additional point is that while dedicated GIS programs handle spatial data by default and display the results in a single way, there are various options in R that must be decided by the user, for example whether to use R's base graphics or a dedicated graphics package such as ggplot2. On the other hand, the main benefits of R for spatial data visualisation lie in the *reproducibility* of its outputs, a feature that we will be using to great effect in this chapter.

R and reproducible research

Finally, there is a drive towards transparency in data and methods datasets in academic publishing. R encourages truly transparent and reproducible research by enabling anyone with an R installation reproduce results described in previous paper. This process is eased by the RStudio integrated development environment (IDE) that allows 'live' R code and results to be embedded in documents. In fact, this tutorial was written in RStudio and can be recompiled on any computer by downloading the project's GitHub repository.

Getting started with the tutorial

The first stage with this tutorial is to download the data from GitHub, where an updated version is stored: github.com/geocomPP/sdvwR. Click on the "Download ZIP" button on the right, and unpack the folder to a sensible place on your computer (e.g. the Desktop). Explore the folder and try opening some of the files, especially those from the sub-folder entitled "data": these are the input datasets we'll be using.

Now open up a version of RStudio on your own computer (or on-line) and we are ready to go.

R and Spatial Data

Preliminaries

R has a unique syntax that is worth learning in basic terms before loading spatial data: to R spatial and non-spatial data are treated in the same way, although they have different underlying data structures. Try typing and running (by pressing `ctl-Enter` in a an RStudio script) the following calculations to see how R works and plot the result.

```
t <- seq(from = 0, to = 20, by = 0.1)
x <- sin(t) * exp(-0.2 * t)
plot(x)
```

Figure 1: unnamed-chunk-1

R code consists of *functions*, usually proceeded by brackets (e.g. `seq`) and *objects* (`d`, `t` and `x`). Each function contains *arguments*, the names of which often do not need to be states: the function `seq(0, 20, 0.1)`, for example, would also work, `from`, `to` and `by` are the *default* arguments. Knowing this is important as it can save typing. In this tutorial, however, we generally spell out each of the argument names, for clarity.

Note the use of the <- assignment to create new objects. Objects are entities that can be called to by name in R and can be renamed through additional assignements (e.g `y <- x` if y seems a more appropriate name).

Spatial Data in R

In any data analysis project, spatial or otherwise, it is important to have a strong understanding of the dataset before progressing. This section will therefore begin with a description of the input data. We will see how data can be loaded into R and exported to other formats, before going into more detail about the underlying structure of spatial data in R.

Loading spatial data in R

In most situations, the starting point of spatial analysis tasks is loading in datasets. These may originate from government agencies, remote sensing devices or 'volunteered geographical information' (Goodchild 2007). R is able to import a very wide range of spatial data formats thanks to its interface with the Geospatial Data Abstraction Library (GDAL), which is enabled by the package `rgdal`. Below we will load data from two spatial data formats: GPS eXchange (`.gpx`) and an ESRI Shapefile.

`readOGR` is in fact capable of loading dozens more file formats, so the focus is on the *method* rather than the specific formats. The `readOGR` function is therefore capable of loading most common spatial file formats. Let's start with a `.gpx` file, a tracklog recording a bicycle ride from Sheffield to Wakefield which was uploaded OpenStreetMap [3].

```
# download.file('http://www.openstreetmap.org/trace/1619756/data', destfile
# = 'data/gps-trace.gpx')
library(rgdal)  # load the gdal package
shf2lds <- readOGR(dsn = "data/gps-trace.gpx", layer = "tracks")  # load track
plot(shf2lds)
shf2lds.p <- readOGR(dsn = "data/gps-trace.gpx", layer = "track_points")  # load points
points(shf2lds.p[seq(1, 3000, 100), ])
```

3

Figure 2: unnamed-chunk-2

In the code above we first used R to *download* a file from the internet, using the function `download.file`. The two essential arguments of this function are `url` (we could have typed `url =` before the link) and `destfile`, the destination file. As with any function, more optional arguments can be viewed by by typing `?download.file`.

When `rgdal` has successfully loaded, the next task is not to import the file directly, but to find out which *layers* are available to import, with `ogrListLayers`. The output from this command tells us that various layers are available, including `tracks` and `track_points`. These are imported into R's *workspace* using `readOGR`.

Finally, the basic `plot` function is used to visualize the newly imported objects, ensuring they make sense. In the second `plot` function, we take a subset of the object (see section . . . for more on this). To see how to add axes, enter `?axis`

Try discovering more about the function by typing `?readOGR`. The documentation explains that the `dsn =` argument is interpreted differently depending on the type of file used. In the above example, the `dsn` was set to as the name of the file. To load Shapefiles, by contrast, the *folder* containing the data is used:

```
lnd <- readOGR(dsn = "data/", "london_sport")
```

Here, the files reside in a folder entitled `data` which in R's current working directory (remember to check this using `getwd()`). If the files were stored in the working directory, one would use `dsn = "."` instead. Again, it may be wise to plot the data that results, to ensure that it has worked correctly. Now that the data has been loaded into R's own `sp` format, try interrogating and plotting it, using functions such as `summary` and `plot`.

The size of spatial datasets in R

Any datasets that have been read into R's *workspace*, which constitutes all objects that can be accessed by name and can be listed using the `ls()` function, can be saved in R's own data storage file type, `.RData`. Spatial datasets can get quite large and this can cause problems on computers by consuming all available memory (RAM) or hard disk space. It is therefore wise to understand roughly how large spatial objects are, providing insight into how long certain functions will take to run.

In the absence of prior knowledge, which of the two objects loaded in the previous section would one expect to be larger? One could hypothesize that the London boroughs represented by the object `lnd` would be larger based on its greater spatial extent, but how much larger? The answer in R is found in the function `object.size`:

```
object.size(shf2lds)
```

```
## 103168 bytes
```

```
object.size(lnd)
```

```
## 79168 bytes
```

Surprisingly, the GPS data is larger. To see why, we can find out how many *vertices* (points connected by lines) are contained in each dataset:

```
shf2lds.f <- fortify(shf2lds)
nrow(shf2lds.f)
```

```
## [1] 6085
```

```
lnd.f <- fortify(lnd)
```

```
## Regions defined for each Polygons
```

```
nrow(lnd.f)
```

```
## [1] 1102
```

In the above block of code we performed two functions for each object: 1) *flatten* the dataset so that each vertice is allocated a unique row 2) use `nrow` to count the result.

It is clear that the GPS data has almost 6 times the number of vertices as does the London data, explaining its larger size. Yet when plotted, the GPS data does not seem more detailed, implying that some of the vertices in the object are not needed for visualisation at the scale of the objects *bounding box*.

Simplifying geometries

In many cases the spatial data we have are too detailed for effective data visualisation. Simplifcation can help to make a graphic more readable and less cluttered. Within the 'rgeos' package it is possible to use the `gSimplify` function to simplify spatial R objects:

```
library(rgeos)
shf2lds.simple <- gSimplify(shf2lds, tol = 0.001)
(object.size(shf2lds.simple)/object.size(shf2lds))[1]
```

```
## [1] 0.03047
```

```
plot(shf2lds.simple)
plot(shf2lds, col = "red", add = T)
```

In the above block of code, `gSimplify` is given the object `shf2lds` and the `tol` argument of 0.001 (much larger tolerance values may be needed, for data that is *projected*). Next, we divide the size of the simplified object by the original (note the use of the / symbol). The output of `0.03...` tells us that the new object is only 3% of its original size. We can see how this has happened by again counting the number of vertices. This time we use the `coordinates` and `nrow` functions together:

```
nrow(coordinates(shf2lds.simple)[[1]][[1]])
```

```
## [1] 44
```

The syntax of the double square brackets will seem strange, providing a taster of how R 'sees' spatial data (see section x). Do not worry about this for now. Of interest here is that the number of vertices has shrunk, from 6,084 to only 44, without loosing much information about the shape of the line. To test this,try plotting the original and simplified tracks on your computer: when visualized using the `plot` function, object `shf2lds.simple` retains the overall shape of the line and is virtually indistinguishable from the original object.

This example is rather contrived because even the larger object `shf2lds` is only 0.107 Mb, negligible compared with the gigabytes of RAM available to modern computers. However, it underlines a wider point: for visualizing *small scale* maps, spatial data *geometries* can often be simplified to reduce processing time and use of computer memory.

Saving and exporting spatial objects

A typical R workflow involves loading the data, processing and finally exporting the data in a new form. `writeOGR`, the logical counterpart of `readOGR` is ideal for this task. Imagine that we want to view the simplified `gpx` data in software that can only read Shapefiles. This is performed using the following command:

```
shf2lds.simple <- SpatialLinesDataFrame(shf2lds.simple, data.frame(row.names = "0",
    a = 1))
writeOGR(shf2lds.simple, layer = "shf2lds", dsn = "data/", driver = "ESRI Shapefile")
```

In the above code, the object was first converted into a spatial dataframe class required by the `writeOGR` command, before being exported as a shapefile entitled shf2lds. Unlike with `readOGR`, the driver must be specified, in this case with "ESRI Shapefile" [4]. The simplified GPS data is now available to other GIS programs for further analysis. Alternatively, `save(shf2lds.simple, file = "data/shf2lds.RData")` will save the object in R's own spatial data format, which is described in the next section.

The structure of spatial data in R

Spatial datasets in R are saved in their own format, defined as `Spatial...` classes within the `sp` package. For this reason, `sp` is the basic spatial package in R, upon which the others depend. Spatial classes range from the basic `Spatial` class to the complex, `SpatialPolygonsDataFrame`: the `Spatial` class contains only two required *slots*[5]:

```
getSlots("Spatial")
```

```
##        bbox proj4string
##    "matrix"       "CRS"
```

This tells us that `Spatial` objects must contain a bounding box (`bbox`) and and CRS. Further details on these can be found by typing `?bbox` and `?proj4string`. All other spatial classes in R build on this foundation of a bounding box and a projection system (which is set automatically to `NA` if it is not known). However, more complex classes contain more slots, some of which are lists which contain additional lists. To find out the slots of `shf2lds.simple`, for example, we would first ascertain its class and then use the `getSlots` command:

```
class(shf2lds.simple)  # identify the object's class
```

```
## [1] "SpatialLinesDataFrame"
## attr(,"package")
## [1] "sp"
```

```
getSlots("SpatialLinesDataFrame")  # find the associated slots
```

```
##         data        lines         bbox  proj4string
## "data.frame"       "list"     "matrix"        "CRS"
```

The same principles apply to all spatial classes including `Spatial* Points`, `Polygons Grids` and `Pixels` as well as associated `*DataFrame` classes. For more information on this, see the `sp` documentation: `?Spatial`.

Manipulating spatial data

Coordinate reference systems

As mentioned in the previous section, all `Spatial` objects in R are allocated a coordinate reference system (CRS). The CRS of any spatial object can be found using the command `proj4string`. In some cases the CRS is not known: in this case the result will simply be `NA`. To discover the CRS of the `lnd` object for example, type the following:

```
proj4string(lnd)
```

```
## [1] "+proj=tmerc +lat_0=49 +lon_0=-2 +k=0.9996012717 +x_0=400000 +y_0=-100000 +ellps=airy "
```

The output may seem cryptic but is in fact highly informative: `lnd` has *projected* coordinates, based on the *Transverse Mercator* system (hence `"+proj=tmerc"` in the output) and its origin is at latitude 49N, -2E. This point is

If we *know* that the CRS is incorrectly specified, it can be re-set. In this case, for example we know that `lnd` actually has a CRS OSGB1936. Knowing also that the code for this is 27700, it can be updated as follows:

```
proj4string(lnd) <- CRS("+init=epsg:27700")
proj4string(lnd)
```

```
## [1] "+init=epsg:27700 +proj=tmerc +lat_0=49 +lon_0=-2 +k=0.9996012717 +x_0=400000 +y_0=-100000 +datum=0:
```

The CRS has now been updated - note that the key details are all the same as before. Note: this method **should never** be used as an attempt to *reproject* data from one CRS to another.

Reprojecting data

Transforming the coordinates of spatial data from one CRS to another (reprojection) is a common task in GIS. This is because data from national sources are generally provided in *projected* coordinates (the location on the cartesian coordinates of a map) whereas data from GPSs and the internet are generally provided in *geographic* coordinates, with latitude and longitude measured in degrees to locate points on the surface of the globe.

Reprojecting data in R is quite simple: all you need is a spatial object with a known CRS and knowledge of the CRS you wish to transform it to. To illustrate why that is necessary, try to plot the objects `lnd` and `shf2lnd.simple` on the same map:

```
combined <- rbind(fortify(shf2lds.simple)[, 1:2], fortify(lnd)[, 1:2])
plot(combined)
```

In the above code we first extracted the coordinates of each point using `fortify` and then plotted them using `plot`. The image shows why reprojection is necessary: the .gpx data are on a totally different scale than the shapefile of London. Hence the tiny dot at the bottom right of the graph. We will now reproject the data, allowing `lnd` and `shf2lds.simple` to be usefully plotted on the same graphic:

```
lnd.wgs84 <- spTransform(lnd, CRSobj = CRS("+init=epsg:4326"))
```

The above code created a new object,`lnd.wgs84`, that contains the same geometries as the original but in a new CRS using the `spTransform` function. The `CRS` argument was set to `"+init=epsg:4326"`, which represents the WGS84 CRS via an EPSG code [6]. Now `lnd` has been reprojected we can plot it next to the GPS data:

```
combined <- rbind(fortify(shf2lds.simple)[, 1:2], fortify(lnd.wgs84)[, 1:2])
plot(combined)
```

Figure 3: Plot of spatial objects with different CRS

Figure 4: Plot of spatial objects sharing the same CRS

Although the plot of the reprojected data is squashed because the axis scales are not fixed and distorted (*geographic* coordinates such as WGS84 should not usually be used for plotting), at least the relative position and shape of both objects can now be seen (making visualisations attractive is covered in the next section). The presence of the dotted line in the top left of the plot confirms our assumption that the GPS data is from around Sheffield, which is northwest of London.

Attribute joins

Because boroughs are official administrative zones, there is much data available at this level that we can link to the polygons in the `lnd` object. We will use the example of crime data to illustrate this data availability, which is stored in the `data` folder available from this project's github page.

```
load("data/crimeAg.Rdata")  # load the crime dataset from an R dataset
```

After the dataset has been explored (e.g. using the `summary` and `head` functions) to ensure compatibility, it can be joined to `lnd`. We will use the the `join` function in the `plyr` package but the `merge` function could equally be used (remember to type `library(plyr)` if needed).

`join` requires all joining variables to have the same name, but this work has already been done [7]. Once this preparation has been done, the join funtion is actually very simple:

```
lnd@data <- join(lnd@data, crimeAg)
```

Take a look at the `lnd@data` object. You should see new variables added, meaning the attribute join was successful.

Spatial joins

A spatial join, like attribute joins, is used to transfer information from one dataset to another. There is a clearly defined direction to spatial joins, with the *target layer* receiving information from another spatial layer based on the proximity of elements from both layers to each other. There are three broad types of spatial join: one-to-one, many-to-one and one-to-many. We will focus only the former two as the third type is rarely used.

One-to-one spatial joins

One-to-one spatial joins are by far the easiest to understand and compute because they simply involve the transfer of attributes in one layer to another, based on location. A one-to-one join is depicted in figure 5 below, and can performed using the same technique as described in the section on spatial aggregation.

Many-to-one spatial joins

Many-to-one spatial joins involve taking a spatial layer with many elements and allocating the attributes associated with these elements to relatively few elements in the target spatial layer. A common type of many-to-one spatial join is the allocation of data collected at many point sources unevenly scattered over space to polygons representing administrative boundaries, as represented in Fig. x.

```
lnd.stations <- readOGR("data/", "lnd-stns", p4s = "+init=epsg:27700")
```

```
## OGR data source with driver: ESRI Shapefile
## Source: "data/", layer: "lnd-stns"
## with 2532 features and 6 fields
## Feature type: wkbPoint with 2 dimensions
```

Figure 5: Illustration of a one-to-one spatial join

Figure 6: Input data for a spatial join

```
plot(lnd)
plot(lnd.stations[round(runif(500, 1, nrow(lnd.stations))), ], add = T)
```

The above code reads in a `SpatialPointsDataFrame` consisting of 2532 transport nodes in and surrounding London and then plots a random sample of 500 of these over the previously loaded borough level administrative boundaries. The reason for piloting a sample of the points rather than all of them is that the boundary data becomes difficult to see if all of the points are piloted. It is also useful to see and practice sampling techniques in practice; try to plot only the first 500 points, rather than a random selection, and describe the difference.

The most obvious issue with the point data from the perspective of a spatial join with the borough data is that many of the points in the dataset are in fact located outside the region of interest. Thus, the first stage in the analysis is to filter the point data such that only those that lie within London's administrative zones are selected. This in itself is a kind of spatial join, and can be accomplished with the following code.

```
proj4string(lnd) <- proj4string(lnd.stations)
lnd.stations <- lnd.stations[lnd, ]   # select only points within lnd
plot(lnd.stations)   # check the result
```


Figure 7: A spatial subset of the points

The station points now clearly follow the form of the `lnd` shape, indicating that the procedure worked. Let's review the code that allowed this to happen: the first line ensured that the CRS associated with each layer is *exactly* the same: this step should not be required in most cases, but it is worth knowing about. Of course, if the coordinate systems are *actually* different in each layer, the function `spTransform` will be needed to make them compatible. This procedure is discussed in section !!!. In this case, only the name was slightly different hence direct alteration of the CRS name via the function `proj4string`.

The second line of code is where the magic happens and the brilliance of R's sp package becomes clear: all that was needed was to place another spatial object in the row index of the points (`[lnd, ]`) and R automatically understood that a subset based on location should be produced. This line of code is an example of R's 'terseness' - only a single line of code is needed to perform what is in fact quite a complex operation.

Spatial aggregation

Now that only stations which *intersect* with the `lnd` polygon have been selected, the next stage is to extract information about the points within each zone. This many-to-one spatial join is also known as *spatial aggregation*.

To do this there are a couple of approaches: one using the `sp` package and the other using `rgeos` (see Bivand et al. 2013, 5.3).

As with the *spatial subset* method described above, the developers of R have been very clever in their implementation of spatial aggregations methods. To minimise typing and ensure consistency with R's base functions, `sp` extends the capabilities of the `aggregate` function to automatically detect whether the user is asking for a spatial or a non-spatial aggregation (they are, in essence, the same thing - we recommend learning about the non-spatial use of `aggregate` in R for comparison).

Continuing with the example of station points in London polygons, let us use the spatial extension of `aggregate` to count how many points are in each borough:

```
lndStC <- aggregate(lnd.stations, by = lnd, FUN = length)
summary(lndStC)
plot(lndStC)
```

As with the spatial subset function, the above code is extremely terse. The aggregate function here does three things: 1) identifies which stations are in which London borough; 2) uses this information to perform a function on the output, in this case `length`, which simply means "count" in this context; and 3) creates a new spatial object equivalent to `lnd` but with updated attribute data to reflect the results of the spatial aggregation. The results, with a legend and colours added, are presented in figure 8 below.


Figure 8: Number of stations in London boroughs

As with any spatial attribute data stored as an `sp` object, we can look at the attributes of the point data using the `@` symbol:

```
head(lnd.stations@data, n = 2)
```

```
##      CODE           LEGEND FILE_NAME NUMBER                 NAME MICE
## 91 5520 Railway Station  gb_south  17607       Belmont Station   19
## 92 5520 Railway Station  gb_south  17608 Woodmansterne Station    5
```

In this case we have three potentially interesting variables: "LEGEND", telling us what the point is, "NAME", and "MICE", which represents the number of mice sightings reported by the public at that point (this is a fictional variable). To illustrate the power of the `aggregate` function, let us use it to find the average number of mice spotted in transport points in each London borough, and the standard deviation:

```
lndAvMice <- aggregate(lnd.stations["MICE"], by = lnd, FUN = mean)
summary(lndAvMice)
lndSdMice <- aggregate(lnd.stations["MICE"], by = lnd, FUN = sd)
summary(lndSdMice)
```

In the above code, `aggregate` was used to create entirely new spatial objects that are exactly the same as `lnd`, except with new attribute data. To add the mean mice count to the original object, the following code can be used:

```
lnd$av.mice <- lndAvMice$MICE
```

The above code creates a new variable in the `lnd@data` object entitled "av.mice" and populates it with desired values. Thus `Spatial` objects can behave in the same way as data.frames when refering to attribute variables.

Summary

To summarise this section, we have taken a look inside R's representation of spatial data, learned how to manipulate these datasets in terms of CRS transformations and attribute data and finally explored spatial joins and aggregation. Only *after* the datasets are well understood and saved in a suitable form should we move on to visualisation.

Fundamentals of Spatial Data Visualisation

Good maps depend on sound analysis and data preparation and can have an enormous impact on the understanding and communication of results. It has never been easier to produce a map. The underlying data required are available in unprecedented volumes and the technological capabilities of transforming them into compelling maps and graphics are increasingly sophisticated and straightforward to use. Data and software, however, only offer the starting points of good spatial data visualisation since they need to be refined and calibrated by the researchers seeking to communicate their findings. In this section we will run through the features of a good map. We will then seek to emulate them with R in Section XX. It is worth noting that not all good maps and graphics contain all the features below – they should simply be seen as suggestions rather than firm principles.

Effective map making is hard process – as Krygier and Wood (XXX) put it "there is a lot to see, think about, and do" (p6). It often comes at the end of a period of intense data analysis and perhaps when the priority is to get a paper finished or results published and can therefore be rushed as a result. The beauty of R (and other scripting languages) is the ability to save code and simply re-run it with different data. Colours, map adornments and other parameters can therefore be quickly applied so it is well worth creating a template script that adheres to best practice.

We have selected ggplot2 as our package of choice for the bulk of our maps and spatial data visualisations because it has a number of these elements at its core. The "gg" in its slightly odd name stands for "Grammar of Graphics", which is a set of rules developed by Leland Wilkinson (2005) in a book of the same name. Grammar in the context of graphics works in much the same way as it does in language- it provides a structure. The structure is informed by both human perception and also mathematics to ensure that the resulting visualisations are both technically sound and comprehensible. Through creating ggplot2, Hadley Wickham, implemented these rules as well as developing ways in which plots can be built up in layers (see Wickham, 2010). This layering component is especially useful in the context of spatial data since it is conceptually the same as map layers in Geographical Information Systems (GIS).

!!!!Maps with ggplot2 !!!!Adding base maps with ggmap

First load the libraries required for this section:

```
library(rgdal)
library(ggplot2)
library(gridExtra)
```

You will also need create a folder and then set it as your working directory. Assuming your name is `Uname`, and the folder is saved as `sdvwR` in the Desktop in Windows, use the following.

```
setwd("C:/Users/Uname/Desktop/sdvwR")
```

For this section we are going to use a map of the world to demonstrate some of the cartographic principles discussed. A world map is available from the Natural Earth website. The code below will download this and save it to your working directory. It is commented out because the data may already be on your system. Uncomment each new line (by deleting the `#` symbol) if you need to download and extract the data.

```
# download.file(url='http://www.naturalearthdata.com/http//www.naturalearthdata.com/download/110m/cultural
# 'ne_110m_admin_0_countries.zip', 'auto') # download file
# unzip('ne_110m_admin_0_countries.zip', exdir = 'data/') # unzip to data
# folder file.remove('ne_110m_admin_0_countries.zip') # remove zip file
```

Once downloaded we can then load the data into the R console.

```
wrld <- readOGR("data/", "ne_110m_admin_0_countries")
```

```
## OGR data source with driver: ESRI Shapefile
## Source: "data/", layer: "ne_110m_admin_0_countries"
## with 177 features and 63 fields
## Feature type: wkbPolygon with 2 dimensions
```

```
plot(wrld)
```


Figure 9: unnamed-chunk-4

To see the first ten rows of attribute information assocuiated with each of the country boundaries type the following:

```
head(wrld@data)[, 1:5]
```

```
##    scalerank        featurecla labelrank           sovereignt sov_a3
## 0          1 Admin-0 country          3          Afghanistan    AFG
## 1          1 Admin-0 country          3               Angola    AGO
## 2          1 Admin-0 country          6              Albania    ALB
## 3          1 Admin-0 country          4 United Arab Emirates    ARE
## 4          1 Admin-0 country          2            Argentina    ARG
## 5          1 Admin-0 country          6              Armenia    ARM
```

You can see there are a lot of columns associated with this file. Although we will keep all of the them, we are only really interested in the population estimate ("pop_est") field. Before progressing it is is worth reprojecting the data in order that the population data can be seen better. The coordinate reference system of the wrld shapefile is currently WGS84. This is the common latitude and longitude format that all spatial software packages understand. From a cartographic perspective the standard plots of this projection, of the kind produced above, are not suitable since they heavily distort the shapes of those countries further from the equator. Instead the Robinson projection provides a good compromise between areal distortion and shape preservation. We therefore project it as follows.

```
wrld.rob <- spTransform(wrld, CRS("+proj=robin"))
plot(wrld.rob)
```
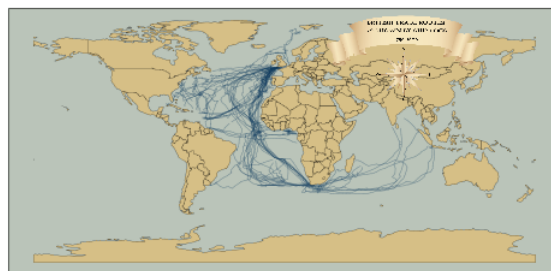

Figure 10: unnamed-chunk-6

"+proj=robin" refers to the Robinson prjection. You will have spotted from the plot that the countries in the world map are much better proportioned.

We now need to "fortify" this spatial data to convert it into a format that ggplot2 understands, we also use "merge" to re-attach the attribute data that is lost in the fortify operation. !!! explain fortify

```
wrld.rob.f <- fortify(wrld.rob, region = "sov_a3")
```

```
## Loading required package: rgeos
## rgeos version: 0.2-19, (SVN revision 394)
##  GEOS runtime version: 3.3.8-CAPI-1.7.8
##  Polygon checking: TRUE
```

```
wrld.pop.f <- merge(wrld.rob.f, wrld.rob@data, by.x = "id", by.y = "sov_a3")
```

15

```
# continuous colour ramp

map <- ggplot(wrld.pop.f, aes(long, lat, group = group, fill = pop_est)) + geom_polygon() +
    coord_equal() + labs(x = "Longitude", y = "Latitude", fill = "World Population") +
    ggtitle("World Population")

# better colours with more breaks- to finish

map + scale_fill_continuous(breaks = )

## Error: argument is missing, with no default


# categorical variables
```

Conforming to colour conventions

Colour has an enormous impact on how people will percieve your graphic. "Readers" of a map come to it with
a range of pre-conceptions about how the world looks. If the map's purpose is to clearly communicate data then
it is often advisable to conform to conventions so as not to disorientate readers to ensure they can focus on the
key messages contained in the data. A good example of this is the use of blue for bodies of water and green for
landmasses. The code example below generates two plots with our wrld.pop.f object. The first colours the land
blue and the sea (in this case the background to the map) green and the second is more conventional. We use the
"grid.arrange" function from the "gridExtra" package to display the maps side by side.

```
map2 <- ggplot(wrld.pop.f, aes(long, lat, group = group)) + coord_equal()

blue <- map2 + geom_polygon(fill = "light blue") + theme(panel.background = element_rect(fill = "dark green

green <- map2 + geom_polygon(fill = "dark green") + theme(panel.background = element_rect(fill = "light bl

grid.arrange(blue, green, ncol = 2)
```

Experimenting with line colour and line widths

In addition to conforming to colour conventions, line colour and width offer important parameters, which are often
overlooked tools for increasing the legibility of a graphic. As the code below demonstrates, it is possible to adjust
line colour through using the "colour" parameter and the line width using the "lwd" parameter. The impact of
different line widths will vary depending on your screen size and resolution. If you save the plot to pdf (or an image)
then the size at which you do this will also affect the line widths.

```
map3 <- map2 + theme(panel.background = element_rect(fill = "light blue"))

yellow <- map3 + geom_polygon(fill = "dark green", colour = "yellow")

black <- map3 + geom_polygon(fill = "dark green", colour = "black")

thin <- map3 + geom_polygon(fill = "dark green", colour = "black", lwd = 0.1)

thick <- map3 + geom_polygon(fill = "dark green", colour = "black", lwd = 1.5)

grid.arrange(yellow, black, thick, thin, ncol = 2)
```

There are other parameters such as layer transparency that can be applied to all aspects of the plot - both points,
lines and polygons - that we will reference in later examples in this chapter.
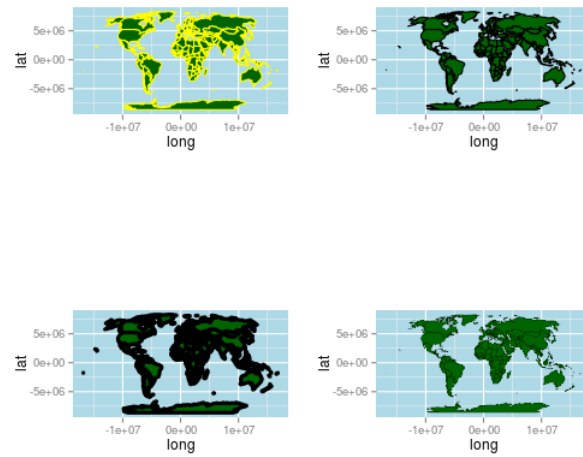
Figure 11: unnamed-chunk-9

Figure 12: unnamed-chunk-10

17

Map Adornments and Annotations

Map adornments and annotations are essential to orientate the viewer and provide context; they include graticules, north arrows, scale bars and data attribution. Not all are required on a single map, indeed it is often best that they are used sparingly to avoid unecessary clutter (Monkhouse and Wilkinson, 1971). Unfortunately it is not always as straightforward to add these in R, and perhaps less so using the ggplot2 paradigm, when compared to a conventional GIS. Here we will outline the ways in which annotations can be added.

!!!! In the maps created so far, we have defined the *aesthetics* of the map in the foundation function `ggplot`. The result of this is that all subsequent layers are expected to have the same variables and essentially contain data with the same dimensions as original dataset. But what if we want to add a new layer from a completely different dataset To do this, we must not add any arguments to the `ggplot` function, only adding data sources one layer at a time:

North arrow

In the maps created so far, we have defined the *aesthetics* of the map in the foundation function `ggplot`. The result of this is that all subsequent layers are expected to have the same variables and essentially contain data with the same dimensions as original dataset. But what if we want to add a new layer from a completely different dataset, e.g. to add an arrow? To do this, we must not add any arguments to the `ggplot` function, only adding data sources one layer at a time:

Here we create an empty plot, meaning that each new layer must be given its own dataset. While more code is needed in this example, it enables much greater flexibility with regards to what can be included in new layer contents. Another possibility is to use the `segment` geom to add a rudimentary arrow (see `?geom_segment` for refinements):

```
library(grid)   # needed for arrow
ggplot() + geom_polygon(data = wrld.pop.f, aes(long, lat, group = group, fill = pop_est)) +
    geom_line(aes(x = c(-1.3e+07, -1.3e+07), y = c(0, 5e+06)), arrow = arrow()) +
    coord_fixed()   # correct aspect ratio
```


Figure 13: unnamed-chunk-11

Scale bar

!!! use geom_segment with geom_text

There is an almost infinite number of different combinations of the above parameters so take inspiration from maps and graphics you have seen and liked. The process is an iterative one, it will take multiple attempts to get right. Show your map to friends and colleagues- all will have an opinion but don't be afraid to stand by the decisions you have taken.

!!!Consistency- across papers.

ggmap is a package that uses the ggplot2 syntax as a template to create maps with image tiles from the likes of Google and OpenStreetMap:

Adding Basemaps To Your Plots

```
library(ggmap)
```

!!! adapt to the lnd.stations object.

The sport object is in British National Grid but the ggmap image tiles are in WGS84. We therefore need to use the lnd.wgs84 object created in the reprojection operation earlier.

The first job is to calculate the bounding box (bb for short) of the lnd.wgs84 object to identify the geographic extent of the image tiles that we need.

```
b <- bbox(lnd.wgs84)
b[1, ] <- (b[1, ] - mean(b[1, ])) * 1.05 + mean(b[1, ])
b[2, ] <- (b[2, ] - mean(b[2, ])) * 1.05 + mean(b[2, ])
# scale longitude and latitude (increase bb by 5% for plot) replace 1.05
# with 1.xx for an xx% increase in the plot size
```

This is then fed into the **get_map** function as the location parameter. The syntax below contains 2 functions. **ggmap** is required to produce the plot and provides the base map data.

```
library(ggmap)
lnd.b1 <- ggmap(get_map(location = b))
```

```
## Warning: bounding box given to google - spatial extent only approximate.
## Warning: cannot open: HTTP status was '403 Forbidden'
```

```
## Error: cannot open URL
## 'http://maps.googleapis.com/maps/api/staticmap?center=51.489304,-0.088233&zoom=10&size=%20640x640&scale=
```

In much the same way as we did above we can then layer the plot with different geoms.

First fortify the lnd.wgs84 object and then merge with the required attribute data (we already did this step to create the sport.f object).

```
lnd.wgs84.f <- fortify(lnd.wgs84, region = "ons_label")
lnd.wgs84.f <- merge(lnd.wgs84.f, lnd.wgs84@data, by.x = "id", by.y = "ons_label")
```

We can now overlay this on our base map.

```
lnd.b1 + geom_polygon(data = lnd.wgs84.f, aes(x = long, y = lat, group = group,
    fill = Partic_Per), alpha = 0.5)
```

The code above contains a lot of parameters. Use the ggplot2 help pages to find out what they are. The resulting map looks okay, but it would be improved with a simpler base map in black and white. A design firm called stamen provide the tiles we need and they can be brought into the plot with the `get_map` function:

```
lnd.b2 <- ggmap(get_map(location = b, source = "stamen", maptype = "toner",
    crop = T))
```

We can then produce the plot as before.

```
lnd.b2 + geom_polygon(data = lnd.wgs84.f, aes(x = long, y = lat, group = group,
    fill = Partic_Per), alpha = 0.5)
```

Finally, if we want to increase the detail of the base map, get_map has a zoom parameter.

```
lnd.b3 <- ggmap(get_map(location = b, source = "stamen", maptype = "toner",
    crop = T, zoom = 11))
```

```
lnd.b3 + geom_polygon(data = lnd.wgs84.f, aes(x = long, y = lat, group = group,
    fill = Partic_Per), alpha = 0.5)
```
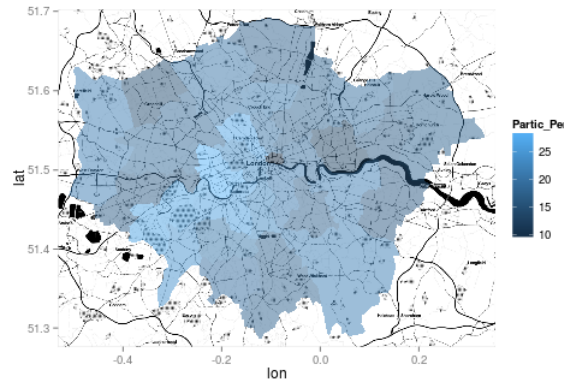

Figure 14: Basemap 3

A Final Example

Here we present a final example that draws upon the many advanced concepts discussed in this chapter to produce a map of 18th Century Shipping flows. The data have been obtained from the CLIWOC project and they represent a sample of digitised ships' logs from the 18th Century. We are using a very small sample of the the full dataset, which is available from here: http://pendientedemigracion.ucm.es/info/cliwoc/. The example has been chosen to demonstrate a range of capabilities within ggplot2 and the ways in which they can be applied to produce high-quality maps with only a few lines of code. We end by showing how the maps can be animated to show the routes over time and the power of R to produce many maps very quickly.

As always, the first step is to load in the required packages and datasets. Here we are using the png package to load in a series of map annotations. These have been created in image editing software and will add a historic feel to the map. We are also loading in a World boundary shapefile and the shipping data itself.

```
library(rgdal)
library(ggplot2)
library(png)
wrld <- readOGR("data/", "ne_110m_admin_0_countries")


## OGR data source with driver: ESRI Shapefile
## Source: "data/", layer: "ne_110m_admin_0_countries"
## with 177 features and 63 fields
## Feature type: wkbPolygon with 2 dimensions


btitle <- readPNG("figure/brit_titles.png")
compass <- readPNG("figure/windrose.png")
bdata <- read.csv("data/british_shipping_example.csv")
```

If you look at the first few lines in the bdata object you will see there are 7 columns with each row representing a single point on the ships course. The year of the journey and the nationality of the ship are also included. The final 3 columns are identifiers that are used later to group the coordinate points together into the paths that ggplot2 plots.

We first specify some plot parameters that remove the axis labels.

```
xquiet <- scale_x_continuous("", breaks = NULL)
yquiet <- scale_y_continuous("", breaks = NULL)
quiet <- list(xquiet, yquiet)
```

The next step is to "fortify" the World coastlines and create the base plot. This sets the extents of the plot window and provides the blank canvas on which we will build up the layers. The first layer created is the wrld object; the code is wrapped in c() to prevent it from executing by simply storing it as the plot's parameters.

```
wrld.f <- fortify(wrld, region = "sov_a3")


## Loading required package: rgeos
## rgeos version: 0.2-19, (SVN revision 394)
##  GEOS runtime version: 3.3.8-CAPI-1.7.8
##  Polygon checking: TRUE


base <- ggplot(wrld.f, aes(x = long, y = lat))
wrld <- c(geom_polygon(aes(group = group), size = 0.1, colour = "black", fill = "#D6BF86",
    data = wrld.f, alpha = 1))
```

To see the result of this simply type:

```
base + wrld
```

The code snipped below creates the plot layer containing the the shipping routes. The geom_path() function is used to string together the coordinates into the routes. You can see within the aes() component we have specified long and lat plus pasted together the trp and group.regroup variables to identify the unique paths.

```
route <- c(geom_path(aes(long, lat, group = paste(bdata$trp, bdata$group.regroup,
    sep = ".")), colour = "#0F3B5F", size = 0.2, data = bdata, alpha = 0.5,
    lineend = "round"))
```

Figure 15: unnamed-chunk-4

We now have all we need to generate the final plot by building the layers together with the + sign. The first 3 arguments are the plot layers and the parameters within theme() are changing the background colour to sea blue. annotation_raster() plots the png map adornments loaded in earlier- this requires the bounding box of each image to be specified. In this case we use latitude and longitude (the plot's coordinate system) and we can use these paramrters to change the png's position and also its size. The final two arguments fix the aspect ratio of the plot and remove the axis labels.

```
base + route + wrld + theme(panel.background = element_rect(fill = "#BAC4B9",
    colour = "black")) + annotation_raster(btitle, xmin = 30, xmax = 140, ymin = 51,
    ymax = 87) + annotation_raster(compass, xmin = 65, xmax = 105, ymin = 25,
    ymax = 65) + coord_equal() + quiet
```

In the plot example we have chosen the colours carefully to give the appearance of a historic map. An alternative approach could be to use a satellite image as a base map. It is possible to use the readPNG function to import NASA's "Blue Marble" image for this purpose. Given that the route information is the same projection as the image - latitude and longitude (WGS84) - it is very straightforward to set the image extent to span -180 to 180 degrees and -90 to 90 degrees and have it align with the shipping data. Producing the plot is accomplished using the code below. This offers a good example of where functionality designed without spatial data in mind can be harnessed for the purposes of producing interesting maps. Once you have produced the plot, alter the code to recolour the shipping routes to make them appear more clearly against the blue marble background.

```
earth <- readPNG("figure/earth_raster.png")

base + annotation_raster(earth, xmin = -180, xmax = 180, ymin = -90, ymax = 90) +
    route + theme(panel.background = element_rect(fill = "#BAC4B9", colour = "black")) +
    annotation_raster(btitle, xmin = 30, xmax = 140, ymin = 51, ymax = 87) +
    annotation_raster(compass, xmin = 65, xmax = 105, ymin = 25, ymax = 65) +
    coord_equal() + quiet
```

Animating your plots

R is not designed to produce animated graphics and as such it has very few functions that enable straightforward animation. To produce animated graphics users can use a loop to plot and then export a series of images that can then be stitched together into a video. There are two approaches to this; the first is to create a loop that fills a

22

Figure 16: unnamed-chunk-6

Figure 17: unnamed-chunk-7

folder with the desired images and then utilise third party software to stitch the images together, whilst the second uses R's own animation package. The latter option still requires the installation of an additional software package called ImageMagick but it has the benefit of creating the animation for you within R and facilitating the export to a range of formats, not least HTML and GIF. Here we demonstrate the use of the package to produce an HTML animation of the shipping tracks completed in each year of the bdata object. The code snippet below appears extremely dense, but it only contains a few addtions to the plot code utilised above.

First load the package:

```r
library(animation)
```

Then clear any previous animation. Obviously the first time you run this it is unnecessary, but it is a good habit to get into.

```r
ani.record(reset = TRUE)
```

We then initiate the "for loop". In this case we are using the unique() function to list the unique years within the bdata object. The loop will take the first year, in this case 1791, and assign it to the object i. The code inside the {} brackets will then run with i=1791. You will spot that i is used in a number of places- first to subset the data when creating the route plot and then as the title in the ggtitle() function. We need to force ggplot to create the graphic within the loop so the entire plot call is wrapped in the print() function. Once the plot is called anin.record() is used to save the plot still and dev.off() used to clear the plot window ready for the next iteration. i is then assigned the next year in the list and the code runs again until all years are plotted.

```r
  for (i in order(unique(bdata$year))
    {
    route<-c(geom_path(aes(long,lat,group=paste(trp, group.regroup, sep=".")), colour="#0F3B5F", size=0.2,
    print(base+route+wrld + theme(panel.background = element_rect(fill='#BAC4B9',colour='black')) +annotat
ani.record()
    dev.off()
    }
```

The final step in the process is to save the animation to HTML and view it in your web browser. ani.replay() retrievesanimation stored by the ani.record() function and outdir=getwd() ensures the final file is stored in your working directory.

```r
saveHTML(ani.replay(), img.name = "record_plot", outdir = getwd())
```

```
## HTML file created at: /home/robin/repos/sdvwR/index.html
```

You will note that there is something a little odd about the order in which the years appear. This can be solved by an additional step before the loop code above. Add this in and then regenerate the animation.

Recap and Conclusions

References

Bivand, R., & Gebhardt, A. (2000). Implementing functions for spatial statistical analysis using the language. Journal of Geographical Systems, 2(3), 307–317.

Bivand, R. S., Pebesma, E. J., & Rubio, V. G. (2008). Applied spatial data: analysis with R. Springer.

Burrough, P. A. & McDonnell, R. A. (1998). Principals of Geographic Information Systems (revised edition). Clarendon Press, Oxford.

Goodchild, M. F. (2007). Citizens as sensors: the world of volunteered geography. GeoJournal, 69(4), 211–221.

Harris, R. (2012). A Short Introduction to R. social-statistics.org.

Kabacoff, R. (2011). R in Action. Manning Publications Co.

Krygier, J. Wood, D. 2011. Making Maps: A Visual Guide to Map Design for GIS (2nd Ed.). New York: The Guildford Press.

Longley, P., Goodchild, M. F., Maguire, D. J., & Rhind, D. W. (2005). Geographic information systems and science. John Wiley & Sons.

Monkhouse, F.J. and Wilkinson, H. R. 1973. Maps and Diagrams Their Compilation and Construction (3rd Edition, reprinted with revisions). London: Methuen & Co Ltd.

Ramsey, P., & Dubovsky, D. (2013). Geospatial Software's Open Future. GeoInformatics, 16(4).

Sherman, G. (2008). Desktop GIS: Mapping the Planet with Open Source Tools. Pragmatic Bookshelf.

Torfs and Brauer (2012). A (very) short Introduction to R. The Comprehensive R Archive Network.

Venables, W. N., Smith, D. M., & Team, R. D. C. (2013). An introduction to R. The Comprehensive R Archive Network (CRAN). Retrieved from http://cran.ma.imperial.ac.uk/doc/manuals/r-devel/R-intro.pdf .

Wickham, H. (2009). ggplot2: elegant graphics for data analysis. Springer.

Wickham, H. 2010. A Layered Grammar of Graphics. American Statistical Association, Institute of Mathematics Statistics and Interface Foundation of North America Journal of Computational and Graphical Statistics. 19, 1: 3-28

Endnotes

1. "What kind of a name is R?" common question. R's name originates from the creators of R, Ross Ihaka and Robert Gentleman. R is an open source implementation of the statistical programming language S, so its name is also a play on words that makes implicit reference to this.

2. R is notoriously difficult to search for on major search engines, as it is such a common letter with many other uses beyond the name of a statistical programming language. This should not be a deterrent, as R has a wealth of excellent online resources. To overcome the issue, you can either be more specific with the search term (e.g. "R spatial statistics") or use an R specific search engine such as rseek.org. You can also search of online help *from within R* using the command `RSiteSearch`. E.g. `RSiteSearch("spatial statistics")`. Experiment and see which you prefer!

3. For more information about this ride, please see robinlovelace.net.

4. A complete list of drivers for importing and exporting spatial data can be displayed by typing `getGDALDriverNames()`.

5. Slots are elements found 'inside' classes of the S4 object system. While the sub-elements of S3 objects such as `data.frame` are referred to using the `$` symbol, the slots of S4 objects are identified using `@`. Thus, the variable x of dataframe `df` can be referred to with `df$x`. In the same way, the data associated with a polygon layer such as `lnd` can be accessed with `lnd@data`. Note that `lnd@data` is itself a dataframe, so can be further specified, e.g. with `lnd@data$name`. For more on spatial data classes, see Bivand et al. (2013).

6. EPSG stands for "European Petroleum Survey Group", but this is not really worth knowing as the organisation is now defunct (www.epsg.org/). The important thing is that EPSG codes provide a unified way to refer to a wide range of coordinate systems, as each CRS has its own epsg code. These can be found at the website spatialreference.org. To see how this website can be useful, try searching for "osgb", for example to find the epsg code for the British National Grid.

7. To see how the `crimeAg` dataset was created, please refer to the "Creating-maps-in-R" tutorial (Cheshire and Lovelace, 2014) hosted on GitHub. The file "intro-spatial-rl.pdf" contains this information, in the section on "Downloading addtional data".

```r
source("chapter.R")  # convert chapter to tex
```