

Introduction

What is R?

R is a free and open source computer program that runs on all major operating systems. R relies primarily on the *command line* for data input: instead of interacting with the program by moving your mouse around clicking on different parts of the screen, users enter commands via the keyboard. This will seem to strange to people accustomed to relying on a graphical user interface (GUI) for most of their computing, e.g. via popular programs such as Microsoft Excel or SPSS, yet the approach has a number of benefits, as highlighted by Gary Sherman (2008, p. 283), developer of the popular GIS program QGIS:

With the advent of “modern” GIS software, most people want to point and click their way through life. That’s good, but there is a tremendous amount of flexibility and power waiting for you with the command line. Many times you can do something on the command line in a fraction of the time you can do it with a GUI.

The joy of this, when you get accustomed to it, is that any command is only ever a few keystrokes away, and the order of the commands sent to R can be stored and repeated in scripts, saving even more time in the long-term (more on this in section ...).

Another important attribute of R, related to its command line interface, is that it is a fully fledged *programming language*. Other GIS programs are written in lower level languages such as C++ which are kept at a safe distance from the users by the GUI. In R, by contrast, the user is ‘close to the metal’ in the sense that what he or she inputs is the same as what R sees when it processes the request. This ‘openness’ can seem raw and daunting to beginners, but it is vital to R’s success. Access to R’s source code and openness about how it works has enabled a veritable army of programmers to improve R over time and add an incredible number of extensions to its base capabilities. Consider for a moment that there are now more than 4000 official packages for R, allowing it to tackle almost any computational or numerical problem one could image, and many more that one could not!

Although writing R source code and creating new packages will not appeal to most R users, it inspires confidence to know that there is a strong and highly skilled community of R developers. If there is a useful spatial function that R cannot currently perform, there is a reasonable chance that someone is working on a solution that will become available at a later date. This constant evolution and improvement is a feature of open source software projects not limited to R, but the range and diversity of extensions is certainly one of its strong points. One area where extension of R’s basic capabilities has been particularly successful is the addition of a wide variety of spatial tools.

The rise of R's spatial capabilities

!!! Quick history of R's spatial packages emphasizing current growth and heavy dependence on sp.

Mention exciting and recently added packages.

Why R for spatial data visualisation?

Aside from confusion surrounding its one character name - “what kind of a name is R?” [1] and “how can you possibly find resources for R online?” [2] - R may also seem a strange choice for a chapter on *spatial* data visualisation specifically. “I thought R was just for statistics?” and “Why not use a proper GIS package like QGIS?” are valid questions.

The first question arises because R was traditionally conceived - and is still primarily known - as a “statistical programming language” (Bivand and Gebhardt 2000). Although R does have cutting edge statistical capabilities, this definition does not do justice to its power and flexibility. Thus, a more accurate albeit longer definition of R is “an integrated suite of software facilities for data manipulation, calculation and graphical display” (Venables et al. 2013). It is important to consider this wider definition before diving into R: it is a fully fledged programming language meaning that it is highly extensible but also that the same result can often be generated in different ways. This can be confusing.

The second question is based on the premise that all ‘proper’ Geographic Information Systems need to operate in the same way, with primacy allocated to a mapping window and a mouse-driven GUI interface. But when we look back at the history of GIS and its definitions, it becomes clear that R *is* fully fledged GIS, when it is set up correctly. All early GIS programs used a command-line interface; GUIs were only developed later as a way to run commands without needing to remember all the command names (although this is largely overcome by good ‘help’ options and auto-completion). A concise definition of a GIS is “a computerized tool for solving geographic problems” (Longley et al. 2005, p. 16) and R certainly enables this. A more expansive definition of GIS is “a powerful set of tools for collecting, storing, retrieving at will, transforming, and displaying spatial data from the real world for a particular set of purposes” (Burrough and McDonnell, 1998, from Bivand et al. 2013, p. 5); R excels at each of these tasks.

That being said, there are a few major differences between R and conventional GIS programs in terms of spatial data visualisation: R is more suited to creating one-off graphics than exploring spatial data through repeated zooming, panning and spatial sub-setting using custom-drawn polygons, compared with conventional GIS programs. Although interactive maps in R can be created (e.g. using the web interface *shiny*), it is recommended that R is used *in addition to* rather than as a direct replacement of dedicated GIS programs, especially now that there are myriad free options to try (Sherman 2008). An additional point is that

while dedicated GIS programs handle spatial data by default and display the results in a single way, there are various options in R that must be decided by the user, for example whether to use R's base graphics or a dedicated graphics package such as `ggplot2`. On the other hand, the main benefits of R for spatial data visualisation lie in the *reproducibility* of its outputs, a feature that we will be using to great effect in this chapter.

R for Reproducible research

!!! Are all the examples going to be reproducible?

All these components - scripting, stability and the ability to embed 'live' code in documents - make R an excellent tool for transparent research. By using R and carefully documenting what has been done, one ensures that the methods used to reach a certain result can be reproduced by anyone anywhere in the world, provided they have access to the input dataset. The RStudio graphical interface with R encourages good documentation. RStudio enables presentations to be created and professional-quality pdf documents to be produced using the custom file formats `.Rpres` and `.rnw`. In fact, this chapter was written in RMarkdown and converted into a pdf document using only RStudio editor!

R in the wild

Examples of where R has had an important visual impact.

Might be good to mention New York Times etc here as key users of R.

An introductory session

The best way to learn to use a new tool is by using it. The metaphor of craft skills is appropriate here: if you wanted to become skilled at scything, for example, you would not spend your time reading about scythes. The same is true of R: the best way to learn how it works is to 'get your hands dirty' and try it out on your own computer. This introductory session will therefore serve as an introduction to R's unique *syntax*, as well an illustration of how other visualisations presented in this chapter can be reproduced.

Chapter overview

including of example dataset used.

Endnotes

1. R's name originates from the creators of R, Ross Ihaka and Robert Gentleman. R is an open source implementation of the statistical programming language S, so its name is also a play on words that makes implicit reference to this.
2. R is notoriously difficult to search for on major search engines, as it is such a common letter with many other uses beyond the name of a statistical programming language. This should not be a deterrent, as R has a wealth of excellent online resources. To overcome the issue, you can either be more specific with the search term (e.g. "R spatial statistics") or use an R specific search engine such as rseek.org. You can also search of online help *from within R* using the command `RSiteSearch`. E.g. `RSiteSearch("spatial statistics")`. Experiment and see which you prefer!

```
# Packages we'll be using for this session
library(rgdal)
library(ggplot2)
library(grid)
# setwd('/Users/james/Dropbox/Abstracts_Papers_Reviews/Geocomp_Chapter')
```

get data- map of the world

To download data of the world the block of code below was used. It is commented out because the data may already be on your system. Uncomment each new line (by deleting the # symbol) if you need to download and extract the data.

```
# download.file(url='http://www.naturalearthdata.com/http://www.naturalearthdata.com/download/
# 'ne_110m_admin_0_countries.zip', 'auto')
# unzip('ne_110m_admin_0_countries.zip', exdir = 'data/')
# file.remove('ne_110m_admin_0_countries.zip')
```

read shape file using rgdal library

```
wrld <- readOGR("data/", "ne_110m_admin_0_countries")

## OGR data source with driver: ESRI Shapefile
## Source: "data/", layer: "ne_110m_admin_0_countries"
## with 177 features and 63 fields
## Feature type: wkbPolygon with 2 dimensions
```

transform this to the robinson projection- this is much better for showing population datasets

```
wrld.rob <- spTransform(wrld, CRS("+init=ESRI:54030"))

## Error: error in evaluating the argument 'CRSobj' in selecting a method for function 'spTransform'

wrld.rob.f <- fortify(wrld, region = "sov_a3")

## Loading required package: rgeos
## rgeos version: 0.3-2, (SVN revision 413M)
## GEOS runtime version: 3.3.9-CAPI-1.7.9
## Polygon checking: TRUE

wrld.pop.f <- merge(wrld.rob.f, wrld@data, by.x = "id", by.y = "sov_a3")
```

continuous colour ramp

```
ggplot(wrld.pop.f, aes(long, lat, group = group, fill = pop_est)) + geom_polygon() +
  coord_equal() + labs(x = "Longitude", y = "Latitude", fill = "World Population") +
  ggtitle("World Population")
```

better colours with more breaks- to finish

```
map+ scale_fill_continuous(breaks=)
```

categorical variables

Conforming to colour conventions

```
map2<- ggplot(wrld.pop.f, aes(long, lat, group = group)) + coord_equal()

blue<-map2+ geom_polygon(fill="light blue") + theme(panel.background =
element_rect(fill = "dark green"))

green<-map2 + geom_polygon(fill="dark green") + theme(panel.background
= element_rect(fill = "light blue"))

grid.arrange(green, blue, ncol=2)
```

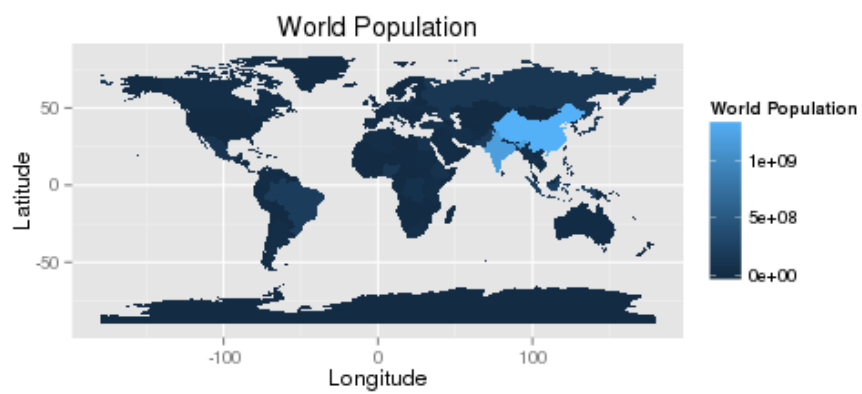


Figure 1: World map with continuous colour ramp fill

Experimenting with line colour and line widths

```
map3<-map2+theme(panel.background = element_rect(fill = "light blue"))
yellow<-map3+ geom_polygon(fill="dark green", colour="yellow")
black<-map3+geom_polygon(fill="dark green", colour="black")
thin<-map3+ geom_polygon(fill="dark green", colour="black", lwd=0.1)
thick<-map3+ geom_polygon(fill="dark green", colour="black", lwd=1.5)
grid.arrange(yellow, black,thick, thin, ncol=2)
```

Annotations

North arrow

In the maps created so far, we have defined the *aesthetics* of the map in the foundation function `ggplot`. The result of this is that all subsequent layers are expected to have the same variables and essentially contain data with the same dimensions as original dataset. But what if we want to add a new layer from a completely different dataset, e.g. to add an arrow? To do this, we must not add any arguments to the `ggplot` function, only adding data sources one layer at a time:

```
arrow <- data.frame(c(2.97, 2.97, 2.9, 3, 3.1, 3.03, 3.03, 2.97), c(1, 4, 4,
  5.5, 4, 4, 1, 1))
names(arrow) <- c("x", "y")
qplot(data = arrow, x = x, y = y)
```

```
qplot(data = arrow, x = x, y = y) + geom_line()
```

```
arrow <- arrow * 5 - 40
```

```
ggplot() + geom_polygon(data = wrld.pop.f, aes(long, lat, group = group, fill = pop_est)) +
  geom_line(data = arrow, aes(x, y))
```

Here we created an empty plot, meaning that each new layer must be given its own dataset. While more code is needed in this example, it enables much greater flexibility with regards to what can be included in new layer contents. Another possibility is to use the `segment` geom to add a pre-made arrow:

```
ggplot() + geom_polygon(data = wrld.pop.f, aes(long, lat, group = group, fill = pop_est)) +
  geom_line(aes(x = c(-160, -160), y = c(0, 25)), arrow = arrow())
```




Figure 2: World map with custom lines added




Figure 3: World map with custom lines added

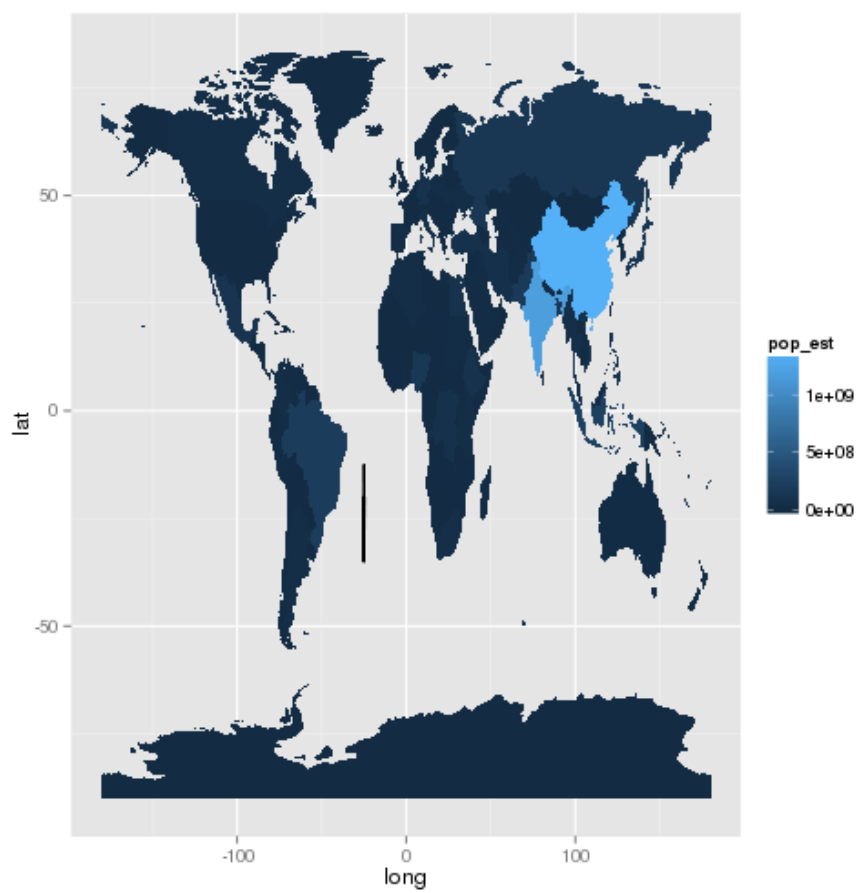


Figure 4: World map with custom lines added

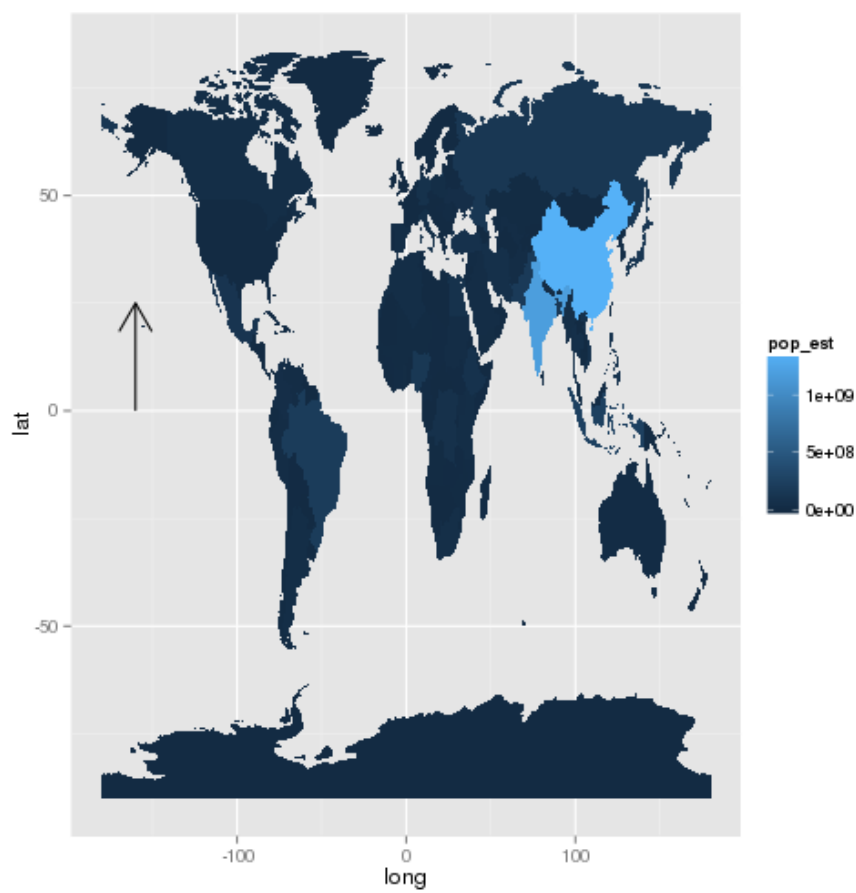


Figure 5: World map with a North arrow

scale bar- found this function

```
hscale_segment = function(breaks, ...) { y = unique(breaksy)stopifnot(length(y) ==
1)dx = max(breaksx) - min(breaksx)dy = 1/30 * dxhscale = data.frame(ix =
min(breaksx), iy=y, jx=max(breaksx), jy = y)vticks = data.frame(ix =
breaksx, iy=(y - dy), jx=breaks$x, jy=(y + dy)) df = rbind(hscale, vticks)
return(geom_segment(data=df, aes(x=ix, xend=jx, y=iy, yend=jy), ...))
}

hscale_text = function(breaks, ...) { dx = max(breaksx) - min(breaksx) dy
= 2/30 * dx breaksy = breaksy + dy return(geom_text(data=breaks, aes(x=x,
y=y, label=label), hjust=0.5, vjust=0, ...))
}
```

R and Spatial Data

Spatial Data in R

In any data analysis project, spatial or otherwise, it is important to have a strong understanding of the dataset before progressing. This section will therefore begin with a description of the input data used in this section. We will see how data can be loaded into R and exported to other formats, before going into more detail about the underlying structure of spatial data in R: how it ‘sees’ spatial data is quite unique.

Loading spatial data in R

In most situations, the starting point of spatial analysis tasks is loading in pre-existing datasets. These may originate from government agencies, remote sensing devices or ‘volunteered geographical information’ from GPS devices, online databases such as Open Street Map or geo-tagged social media (Goodchild 2007). The diversity of geographical data formats is large.

R is able to import a very wide range of spatial data formats thanks to its interface with the Geospatial Data Abstraction Library (GDAL), which is enabled by loading the package `rgdal` into R. Below we will load data from two spatial data formats: GPS eXchange (`.gpx`) and an ESRI Shapefile (consisting of at least files with `.shp`, `.shx` and `.dbf` extensions).

`readOGR` is in fact capable of loading dozens more file formats, so the focus is on the *method* rather than the specific formats. The ‘take home message’ is that the `readOGR` function is capable of loading most common spatial file formats, but behaves differently depending on file type. Let’s start with a `.gpx` file, a

tracklog recording a bicycle ride from Sheffield to Wakefield which was uploaded Open Street Map. [!!! more detail?]

```
# download.file('http://www.openstreetmap.org/trace/1619756/data', destfile
# = 'data/gps-trace.gpx')
library(rgdal) # load the gdal package

## Loading required package: sp
## rgdal: version: 0.8-11, (SVN revision 479M)
## Geospatial Data Abstraction Library extensions to R successfully loaded
## Loaded GDAL runtime: GDAL 1.9.2, released 2012/10/08
## Path to GDAL shared files: /usr/share/gdal
## Loaded PROJ.4 runtime: Rel. 4.8.0, 6 March 2012, [PJ_VERSION: 480]
## Path to PROJ.4 shared files: (autodetected)

ogrListLayers(dsn = "data/gps-trace.gpx") # which layers are available?

## [1] "waypoints"      "routes"          "tracks"          "route_points"
## [5] "track_points"

shf2lds <- readOGR(dsn = "data/gps-trace.gpx", layer = "tracks") # load track

## OGR data source with driver: GPX
## Source: "data/gps-trace.gpx", layer: "tracks"
## with 1 features and 12 fields
## Feature type: wkbMultiLineString with 2 dimensions

plot(shf2lds)
shf2lds.p <- readOGR(dsn = "data/gps-trace.gpx", layer = "track_points") # load points

## OGR data source with driver: GPX
## Source: "data/gps-trace.gpx", layer: "track_points"
## with 6085 features and 26 fields
## Feature type: wkbPoint with 2 dimensions

points(shf2lds.p[seq(1, 3000, 100), ])
```

There is a lot going on in the preceding 7 lines of code, including functions that you are unlikely to have encountered before. Let us think about what has happened, line-by-line.

First, we used R to *download* a file from the internet, using the function `download.file`. The two essential arguments of this function are `url` (we could have typed `url = before the link`) and `destfile` (which means destination



Figure 6: Leeds to Sheffield GPS data

file). As with any function, more optional arguments can be viewed by typing `?download.file`.

When `rgdal` has successfully loaded, the next task is not to import the file directly, but to find out which *layers* are available to import, with the function `ogrListLayers`. The output from this command tells us that various layers are available, including `tracks` and `track_points`, which we subsequently load using `readOGR`. The basic `plot` function is used to plot the newly imported objects, ensuring they make sense. In the second `plot` function, we take a subset of the object (see section ... for more on this).

As stated in the help documentation (accessed by entering `?readOGR`), the `dsn` = argument is interpreted differently depending on the type of file used. In the above example, the filename was the data source name. To load Shapefiles, by contrast, the *folder* containing the data is used:

```
lnd <- readOGR(dsn = "data/", "london_sport")
```

Here, the data is assumed to reside in a folder entitled `data` which in R's current working directory (remember to check this using `getwd()`). If the files were stored in the working directory, one would use `dsn = "."` instead. Again, it may be wise to plot the data that results, to ensure that it has worked correctly. Now that the data has been loaded into R's own `sp` format, try interrogating and plotting it, using functions such as `summary` and `plot`.

The size of spatial datasets in R

Any data that has been read into R's *workspace*, which constitutes all objects that can be accessed by name and can be listed using the `ls()` function, can be saved in R's own data storage file type, `.RData`. Spatial datasets can get quite large and this can cause problems on computers by consuming all available random access memory (RAM) or hard disk space available to the computer. It is therefore wise to understand roughly how large spatial objects are; this will also provide insight into how long certain functions will take to run.

In the absence of prior knowledge, which of the two objects loaded in the previous section would be expected to take up more memory. One could hypothesise that the London boroughs represented by the object `lnd` would be larger, but how much larger? We could simply look at the size of the associated files, but R also provides a function (`object.size`) for discovering how large objects loaded into its workspace are:

```
object.size(shf2lds)
```

```
## 107464 bytes
```

```
## 125544 bytes
```

```
sapply(lnd@polygons, function(x) length(x))
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

```
x <- sapply(lnd@polygons, function(x) nrow(x@Polygons[[1]]@coords))
sum(x)
```

```
## [1] 1102
```

```
sapply(shf2lds@lines, function(x) length(x))
```

```
## [1] 1
```

```
sapply(shf2lds@lines, function(x) nrow(x@Lines[[1]]@coords))
```

```
## [1] 6085
```

It is quite likely that the above code little sense at first; the important thing to remember is that for each object we performed two functions: 1) a check that each line or polygon consists only of a single *part* (that can be joined to attribute data) and 2) the use of `nrow` to count the number of vertices. The use of the `@` symbol should seem strange - its meaning will become clear in the section !!! (Note also that the function `fortify`, discussed in section !!!, can also be used to extract the vertice count of spatial objects in R.)

Without worrying, for now, about how these vertex counts were performed, it is clear that the GPS data has almost 6 times the number of vertices as does the London data, explaining its larger size. Yet when plotted, the GPS data does not seem more detailed, implying that some of the vertices in the object are not needed for visualisation at the scale of the objects *bounding box*.

Simplifying geometries

The wastefulness of the GPS data for visualisation (the full dataset may be useful for other types of analysis) raises the question following question: can the object be simplified such that its key features remain while substantially reducing its size? The answer is yes. In the code below, we harness the power of the `rgeos` package and its `gSimplify` function to simplify spatial R objects (the code can also be used to simplify polygon geometries):

```
library(rgeos)

## rgeos version: 0.3-2, (SVN revision 413M)
## GEOS runtime version: 3.3.9-CAPI-1.7.9
## Polygon checking: TRUE

shf2lds.simple <- gSimplify(shf2lds, tol = 0.001)
(object.size(shf2lds.simple)/object.size(shf2lds))[1]

## [1] 0.04608

plot(shf2lds.simple)
plot(shf2lds, col = "red", add = T)
```

In the above block of code, `gSimplify` is given the object `shf2lds` and the `tol` argument, short for “tolerance”, is set at 0.001 (much larger values may be needed, for data that use is *projected* - does not use latitude and longitude). Comparison between the sizes of the simplified object and the original shows that the new object is less than 3% of its original size. Try plotting the original and simplified tracks on your computer: when visualised using the `plot` function, it becomes clear that the object `shf2lds.simple` retains the overall shape of the line and is virtually indistinguishable from the original object.

This example is rather contrived because even the larger object `shf2lds` is only 0.107 Mb, negligible compared with the gigabytes of RAM available to modern computers. However, it underlines a wider point: for *visualisation* purposes at small spatial scales (i.e. covering a large area of the Earth on a small map), the *geometries* associated with spatial data can often be simplified to reduce processing time and usage of RAM. The other advantage of simplification is that it reduces the size occupied by spatial datasets when they are saved.

Saving and exporting spatial objects

The structure of spatial data in R

Spatial* data

Points

Lines

Polygons

Grids and raster data

‘Flattening’ data with `fortify`

The main spatial packages

`sp`

`rgdal`

`rgeos`

Maps with `ggplot2`

Adding base maps with `ggmap`

Manipulating spatial data

Coordinate reference systems and transformations

Attribute joins

Spatial joins

A spatial join, like attribute joins, is used to transfer information from one dataset to another. There is a clearly defined direction to spatial joins, with the *target layer* receiving information from another spatial layer based on the proximity of elements from both layers to each other. There are three broad types of spatial join: one-to-one, many-to-one and one-to-many. We will focus only the former two as the third type is rarely used.

One-to-one spatial joins are by far the easiest to understand and compute because they simply involve the transfer of attributes in one layer to another, based on location. A one-to-one join is depicted in figure x below.

Many-to-one spatial joins involve taking a spatial layer with many elements and allocating the attributes associated with these elements to relatively few elements in the target spatial layer. A common type of many-to-one spatial join is the allocation of data collected at many point sources unevenly scattered over space to polygons representing administrative boundaries, as represented in Fig. x.



Figure 7: Illustration of a one-to-one spatial join

```

lnd.stations <- readOGR("data/", "lnd-stns", p4s = "+init=epsg:27700")

## OGR data source with driver: ESRI Shapefile
## Source: "data/", layer: "lnd-stns"
## with 2532 features and 6 fields
## Feature type: wkbPoint with 2 dimensions

plot(lnd)
plot(lnd.stations[round(runif(n = 500, min = 1, max = nrow(lnd.stations))),
], add = T)

```




Figure 8: Input data for a spatial join

The above code reads in a `SpatialPointsDataFrame` consisting of 2532 transport nodes in and surrounding London and then plots a random sample of 500 of

these over the previously loaded borough level administrative boundaries. The reason for plotting a sample of the points rather than all of them is that the boundary data becomes difficult to see if all of the points are plotted. It is also useful to see and practice sampling techniques in practice; try to plot only the first 500 points, rather than a random selection, and describe the difference.

The most obvious issue with the point data from the perspective of a spatial join with the borough data is that many of the points in the dataset are in fact located outside the region of interest. Thus, the first stage in the analysis is to filter the point data such that only those that lie within London's administrative zones are selected. This in itself is a kind of spatial join, and can be accomplished with the following code.

```
proj4string(lnd) <- proj4string(lnd.stations)

## Warning: A new CRS was assigned to an object with an existing CRS:
## +proj=tmerc +lat_0=49 +lon_0=-2 +k=0.9996012717 +x_0=400000 +y_0=-100000 +ellps=airy +units=m +no_defs
## without reprojecting.
## For reprojection, use function spTransform in package rgdal

lnd.stations <- lnd.stations[lnd, ] # select only points within lnd
plot(lnd.stations) # check the result
```

The station points now clearly follow the form of the `lnd` shape, indicating that the procedure worked. Let's review the code that allowed this to happen: the first line ensured that the CRS associated with each layer is *exactly* the same: this step should not be required in most cases, but it is worth knowing about. Of course, if the coordinate systems are *actually* different in each layer, the function `spTransform` will be needed to make them compatible. This procedure is discussed in section !!! In this case, only the name was slightly different hence direct alteration of the CRS name via the function `proj4string`.

The second line of code is where the magic happens and the brilliance of R's `sp` package becomes clear: all that was needed was to place another spatial object in the row index of the points (`[lnd,]`) and R automatically understood that a subset based on location should be produced. This line of code is an example of R's 'terseness' - only a single line of code is needed to perform what is in fact quite a complex operation.

Spatial aggregation

Now that only stations which *intersect* with the `lnd` polygon have been selected, the next stage is to extract information about the points within each zone. This many-to-one spatial join is also known as *spatial aggregation*. To do this there




Figure 9: A spatial subset of the points

are a couple of approaches: one using the `sp` package and the other using `rgeos` (see Bivand et al. 2013, 5.3).

As with the *spatial subset* method described above, the developers of R have been very clever in their implementation of spatial aggregations methods. To minimise typing and ensure consistency with R's base functions, `sp` extends the capabilities of the `aggregate` function to automatically detect whether the user is asking for a spatial or a non-spatial aggregation (they are, in essence, the same thing - we recommend learning about the non-spatial use of `aggregate` in R for comparison).

Continuing with the example of station points in London polygons, let us use the spatial extension of `aggregate` to count how many points are in each borough:

```
lndStC <- aggregate(lnd.stations, by = lnd, FUN = length)
summary(lndStC)
plot(lndStC)
```

As with the spatial subset function, the above code is extremely terse. The `aggregate` function here does three things: 1) identifies which stations are in which London borough; 2) uses this information to perform a function on the output, in this case `length`, which simply means “count” in this context; and 3) creates a new spatial object equivalent to `lnd` but with updated attribute data to reflect the results of the spatial aggregation. The results, with a legend and colours added, are presented in Fig !!! below.

As with any spatial attribute data stored as an `sp` object, we can look at the attributes of the point data using the `@` symbol:

```
head(lnd.stations@data)
```

##	CODE	LEGEND	FILE_NAME	NUMBER	NAME	MICE
## 91	5520	Railway Station	gb_south	17607	Belmont Station	19
## 92	5520	Railway Station	gb_south	17608	Woodmansterne Station	5
## 93	5520	Railway Station	gb_south	17609	Coulsdon South Station	11
## 94	5520	Railway Station	gb_south	17610	Smitham Station	14
## 95	5520	Railway Station	gb_south	17611	Kenley Station	11
## 96	5520	Railway Station	gb_south	17612	Reedham Station	8

In this case we have three potentially interesting variables: “LEGEND”, telling us what the point is, “NAME”, and “MICE”, which represents the number of mice sightings reported by the public at that point (this is a fictional variable). To illustrate the power of the `aggregate` function, let us use it to find the average number of mice spotted in transport points in each London borough, and the standard deviation:




Figure 10: Number of stations in London boroughs


```
lndAvMice <- aggregate(lnd.stations["MICE"], by = lnd, FUN = mean)
summary(lndAvMice)
lndSdMice <- aggregate(lnd.stations["MICE"], by = lnd, FUN = sd)
summary(lndSdMice)
```

Clipping

References

- Bivand, R., & Gebhardt, A. (2000). Implementing functions for spatial statistical analysis using the language. *Journal of Geographical Systems*, 2(3), 307–317.
- Bivand, R. S., Pebesma, E. J., & Rubio, V. G. (2008). *Applied spatial data: analysis with R*. Springer.
- Burrough, P. A. & McDonnell, R. A. (1998). *Principals of Geographic Information Systems* (revised edition). Clarendon Press, Oxford.
- Goodchild, M. F. (2007). Citizens as sensors: the world of volunteered geography. *GeoJournal*, 69(4), 211–221.
- Harris, R. (2012). A Short Introduction to R. social-statistics.org.
- Kabacoff, R. (2011). *R in Action*. Manning Publications Co.
- Krygier, J. Wood, D. 2011. *Making Maps: A Visual Guide to Map Design for GIS* (2nd Ed.). New York: The Guildford Press.
- Longley, P., Goodchild, M. F., Maguire, D. J., & Rhind, D. W. (2005). *Geographic information systems and science*. John Wiley & Sons.
- Ramsey, P., & Dubovsky, D. (2013). Geospatial Software’s Open Future. *GeoInformatics*, 16(4).
- Sherman, G. (2008). *Desktop GIS: Mapping the Planet with Open Source Tools*. Pragmatic Bookshelf.
- Torfs and Brauer (2012). A (very) short Introduction to R. The Comprehensive R Archive Network.
- Venables, W. N., Smith, D. M., & Team, R. D. C. (2013). An introduction to R. The Comprehensive R Archive Network (CRAN). Retrieved from <http://cran.ma.imperial.ac.uk/doc/manuals/r-devel/R-intro.pdf>.
- Wickham, H. (2009). *ggplot2: elegant graphics for data analysis*. Springer.